

fork + exec + dup

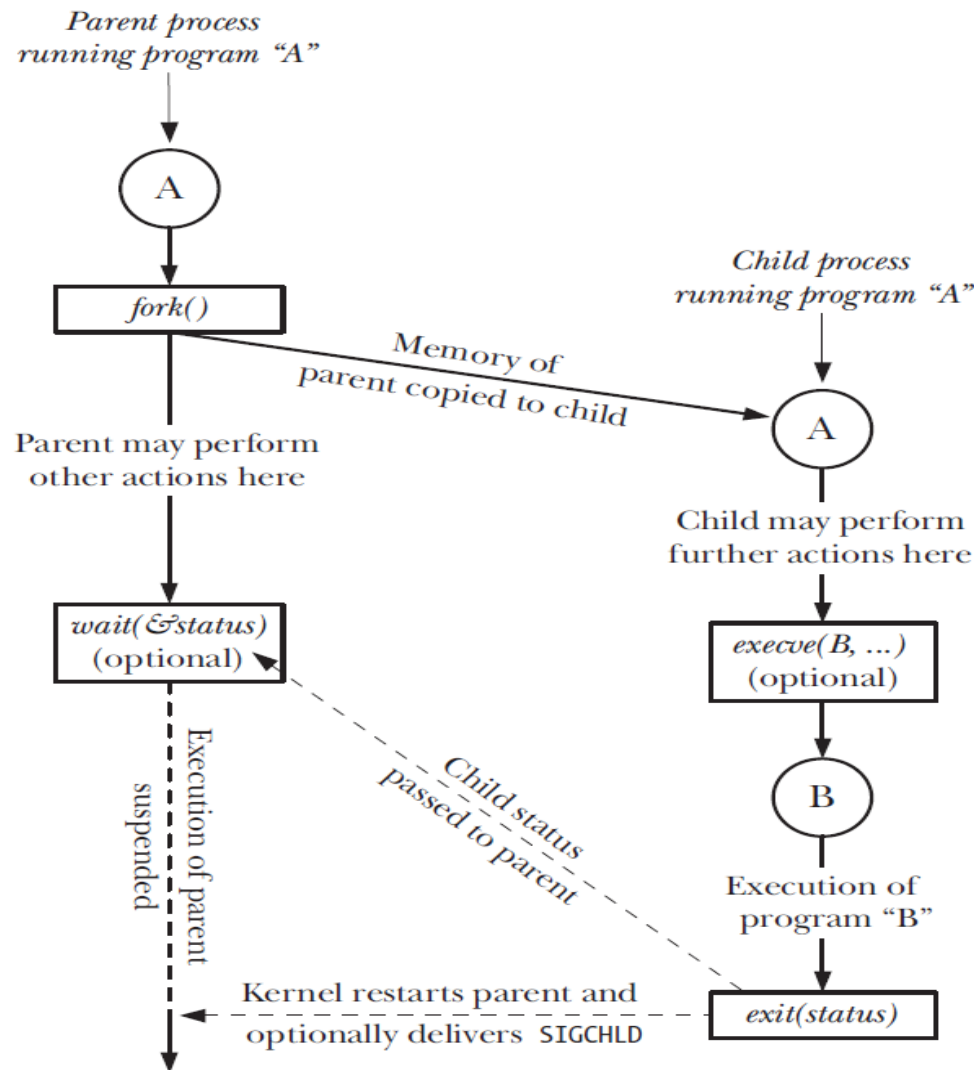
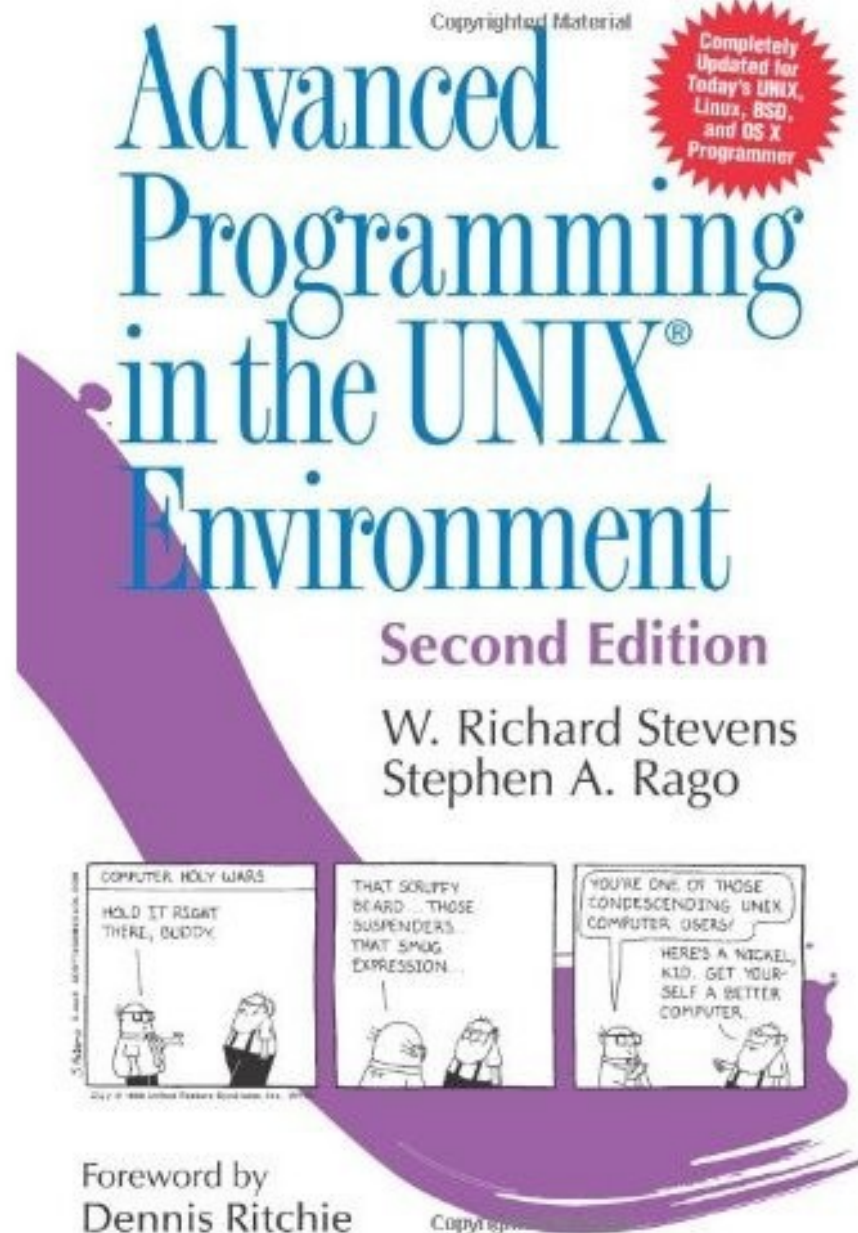


Figure 24-1: Overview of the use of `fork()`, `exit()`, `wait()`, and `execve()`



ADDISON-WESLEY PROFESSIONAL COMPUTING SERIES

Fourth Edition



OPERATING SYSTEM CONCEPTS

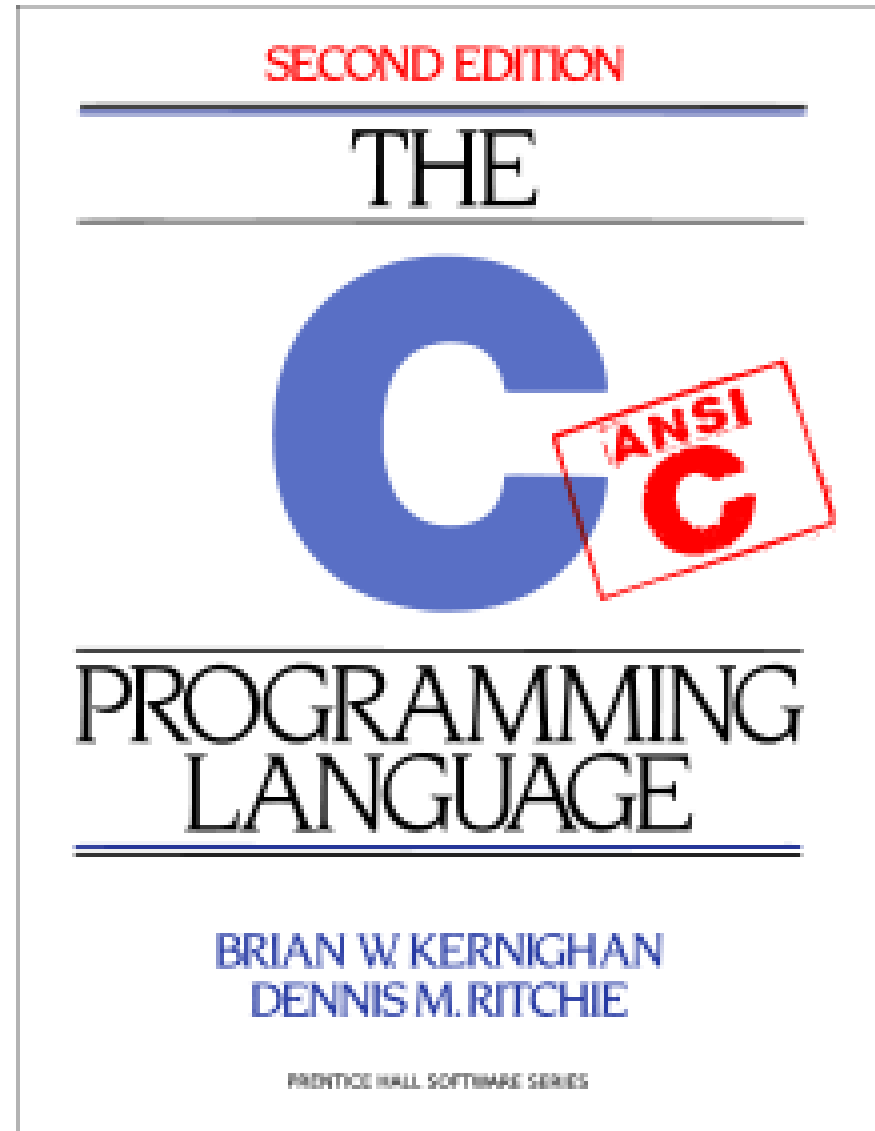
Abraham Silberschatz

Peter Baer Galvin

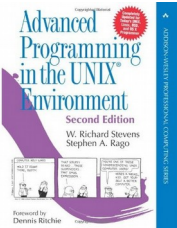
Greg Gagne

Ninth Edition



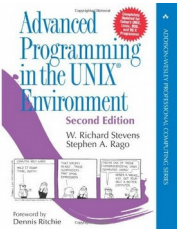


W. Richard Stevens (APUE/2E), Fig. 8.1



```
#include "apue.h"
int      glob = 6;          /* external variable in initialized data */
char     buf[] = "a write to stdout\n";
int
main(void)
{
    int      var;           /* automatic variable on the stack */
    pid_t    pid;
    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */
    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid == 0) {   /* child */
        glob++;             /* modify variables */
        var++;
    } else {
        sleep(2);           /* parent */
    }
    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```

W. Richard Stevens (APUE/2E), Fig. 8.1



```
$ ./a.out
```

```
a write to stdout
```

```
before fork
```

```
pid = 430, glob = 7, var = 89
```

```
pid = 429, glob = 6, var = 88
```

Note 1

Note 2

*child's variables were changed
parent's copy was not changed*

```
$ ./a.out > temp.out
```

```
$ cat temp.out
```

```
a write to stdout
```

```
before fork
```

```
pid = 432, glob = 7, var = 89
```

```
before fork
```

```
pid = 431, glob = 6, var = 88
```

Note 3

Note 4

Note 5

Note 6

Note 7

Notes 1 & 3: the function write is not buffered.

Note 2: standard output (in standard I/O library) is **line** buffered if it's connected to a terminal device. The standard output buffer is flushed by the newline.

Notes 4-7: standard output (redirected to a file) is **fully** buffered. The printf (before the fork) is called once, but the line remains in the buffer when fork is called. This buffer is then copied into the child when the parent's data space is copied to the child. Both the parent and the child now have a standard I/O buffer with this line in it. The second printf (right before the exit), just appends its data to the existing buffer. When each process terminates, its copy of the buffer is finally flushed.

The basic I/O functions `open`, `read`, `write`, `close` are called **unbuffered** I/O functions because **each** read or write invokes a system call into the kernel.

The goal of the buffering provided by the standard I/O library is to use the minimum number of read and write calls.

Three types of buffering are provided:

1. **Fully buffered.** In this case, actual I/O takes place when the standard I/O buffer is filled. The term *flush* describes the writing of a standard I/O buffer. A buffer can be flushed automatically by the standard I/O routines, such as when a buffer fills, or we can call the function `fflush` to flush a stream.

2. **Line buffered.** In this case, the standard I/O library performs I/O when a newline character is encountered on input or output.

3. Unbuffered. The standard I/O library does not buffer the characters.

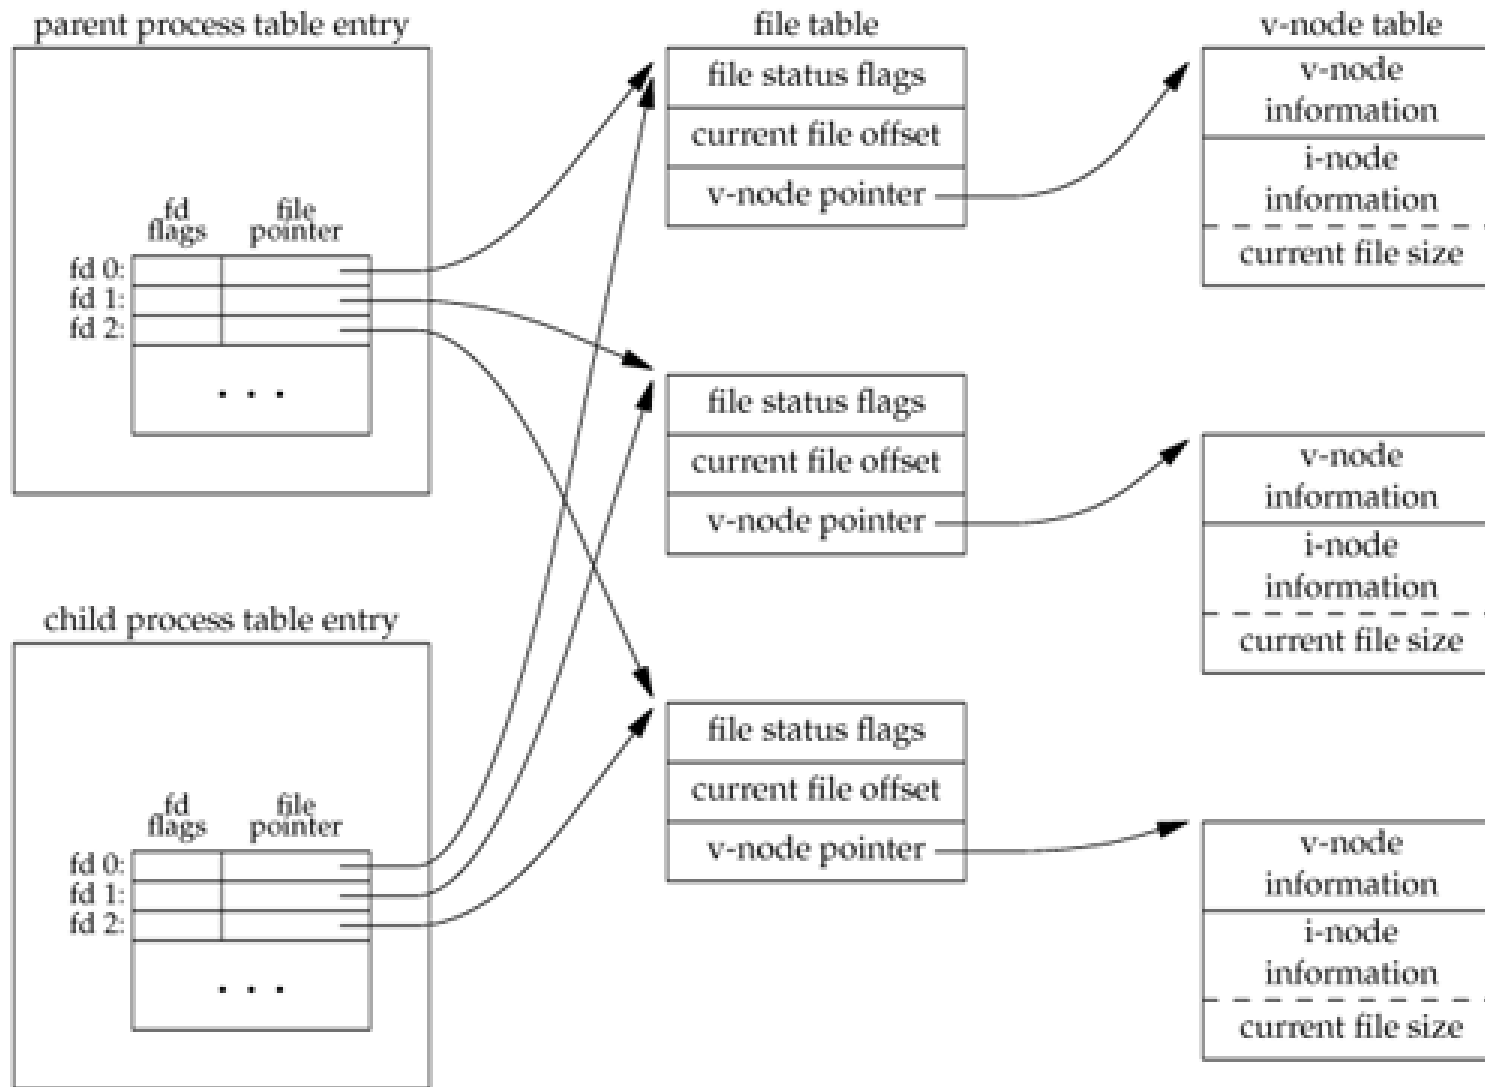
ISO C requires the following buffering characteristics:

- Standard input and standard output are fully buffered, if and only if they do not refer to an interactive device.
- Standard error is never fully buffered.

Most implementations default to the following types of buffering:

- Standard error is always **unbuffered**.
- All other streams are **line buffered** if they refer to a terminal device; otherwise, they are **fully buffered**.

APUE, Fig. 8.2. Sharing of open files between parent and child after fork





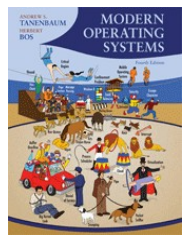
MOS4E, Fig. 1.19. A stripped-down shell

```
#define TRUE 1

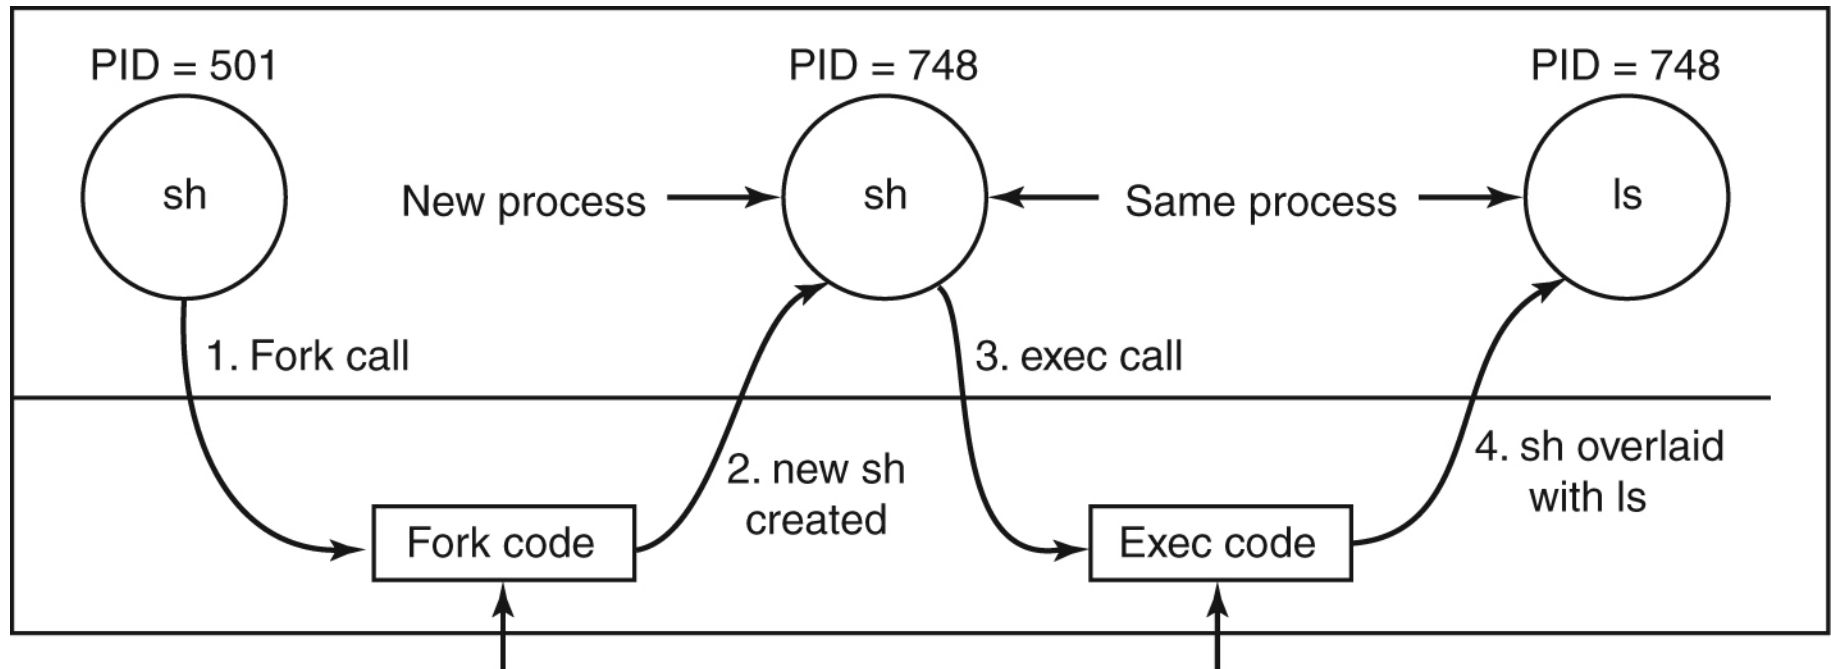
while (TRUE) {
    type_prompt( );
    read_command(command, parameters);

    if (fork( ) != 0) {
        /* Parent code. */
        waitpid(-1, &status, 0);
    } else {
        /* Child code. */
        execve(command, parameters, 0);
    }
}
```

/* repeat forever */
/* display prompt on the screen */
/* read input from terminal */
/* fork off child process */
/* wait for child to exit */
/* execute command */



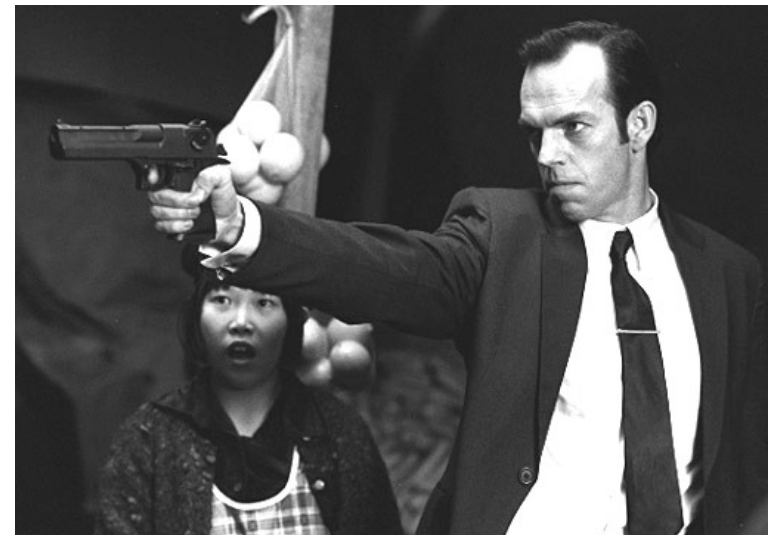
MOS4E, Fig. 10.8. The steps in executing the command `ls` typed to the shell



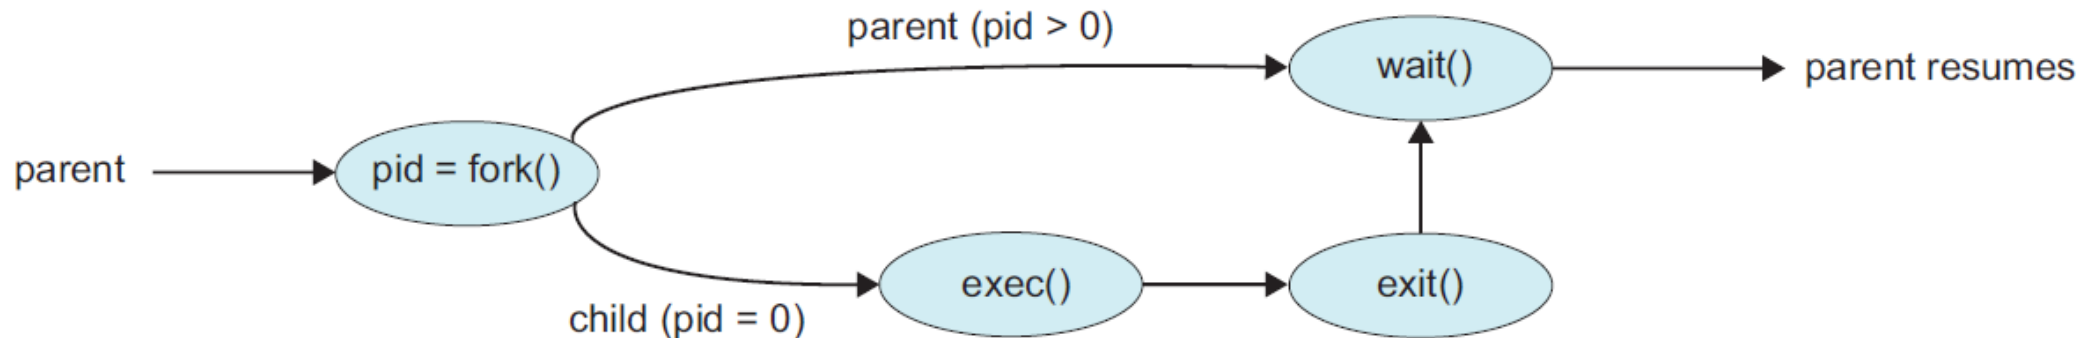
Allocate child's task structure
 Fill child's task structure from parent
 Allocate child's stack and user area
 Fill child's user area from parent
 Allocate PID for child
 Set up child to share parent's text
 Copy page tables for data and stack
 Set up sharing of open files
 Copy parent's registers to child

Find the executable program
 Verify the execute permission
 Read and verify the header
 Copy arguments, environ to kernel
 Free the old address space
 Allocate new address space
 Copy arguments, environ to stack
 Reset signals
 Initialize registers

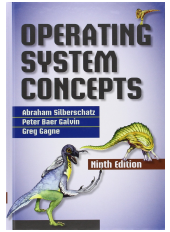
Si quitar la libertad de hacer exec ...



OSC/9E, Fig. 3.10. Process creation using the fork() system call



OSC/9E, Fig. 3.11. Creating a separate process using the Win32 API.



```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL,                /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL,                               /* don't inherit process handle */
        NULL,                               /* don't inherit thread handle */
        FALSE,                             /* disable handle inheritance */
        0,                                  /* no creation flags */
        NULL,
        NULL,
        &pi,
        &si))
    ...
}
```

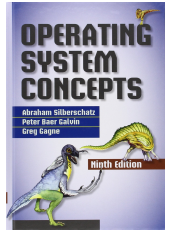
```

...
    NULL,                /* use parent's environment block */
    NULL,                /* use parent's existing directory */
    &si,
    &pi))
{
    fprintf(stderr, "Create Process Failed");
    return -1;
}
/* parent will wait for the child to complete */
WaitForSingleObject(pi.hProcess, INFINITE);
printf("Child Complete");

/* close handles */
CloseHandle(pi.hProcess);
CloseHandle(pi.hThread);
}

```

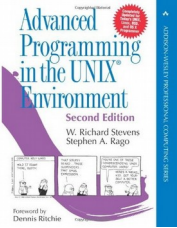
OSC/9E, Fig. 3.9. Process creation



```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t  pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
    }
    return 0;
}
```

W. Richard Stevens (APUE/2E), 8.10 exec Functions



```
#include <unistd.h>
```

```
int execl(const char *pathname, const char *arg0, ... /* (char *)0 */ );
```

```
int execv(const char *pathname, char *const argv []);
```

```
int execle(const char *pathname, const char *arg0, ...  
           /* (char *)0, char *const envp[] */ );
```

```
int execve(const char *pathname, char *const argv[], char *const envp []);
```

```
int execlp(const char *filename, const char *arg0, ... /* (char *)0 */ );
```

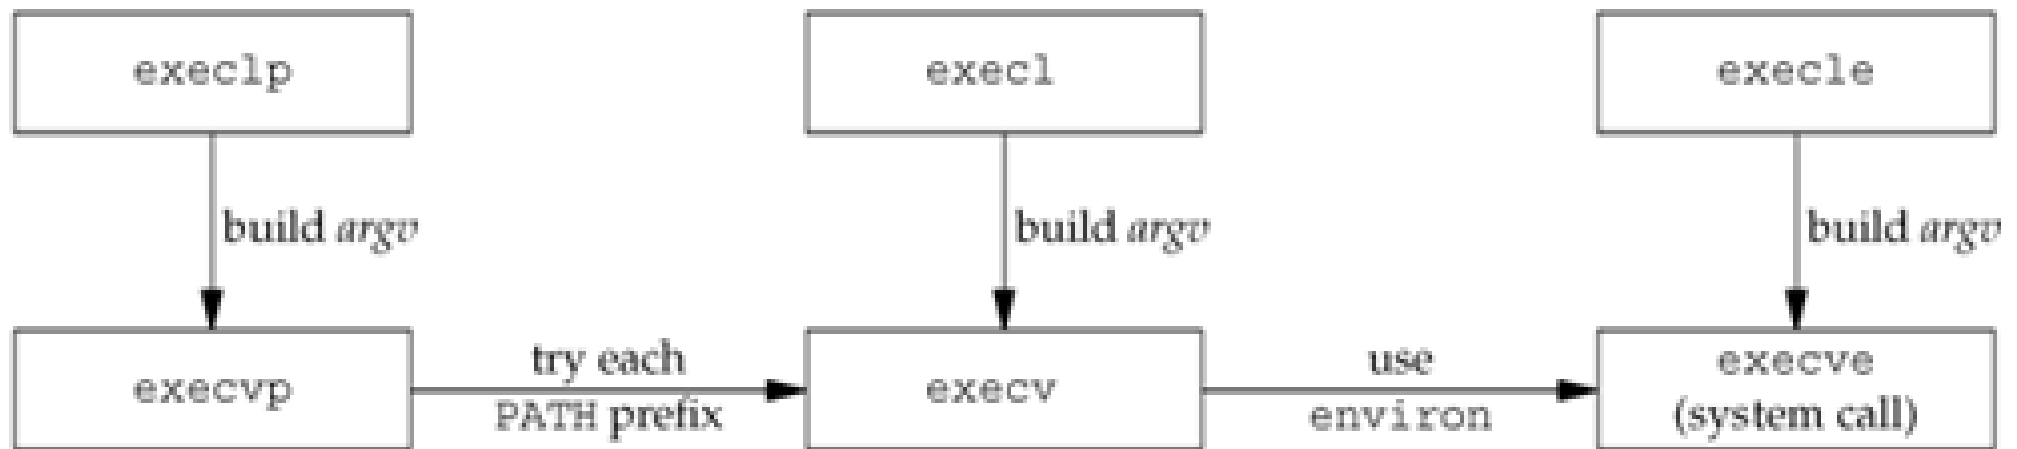
```
int execvp(const char *filename, char *const argv []);
```

All six return: -1 on error, **no return** on success

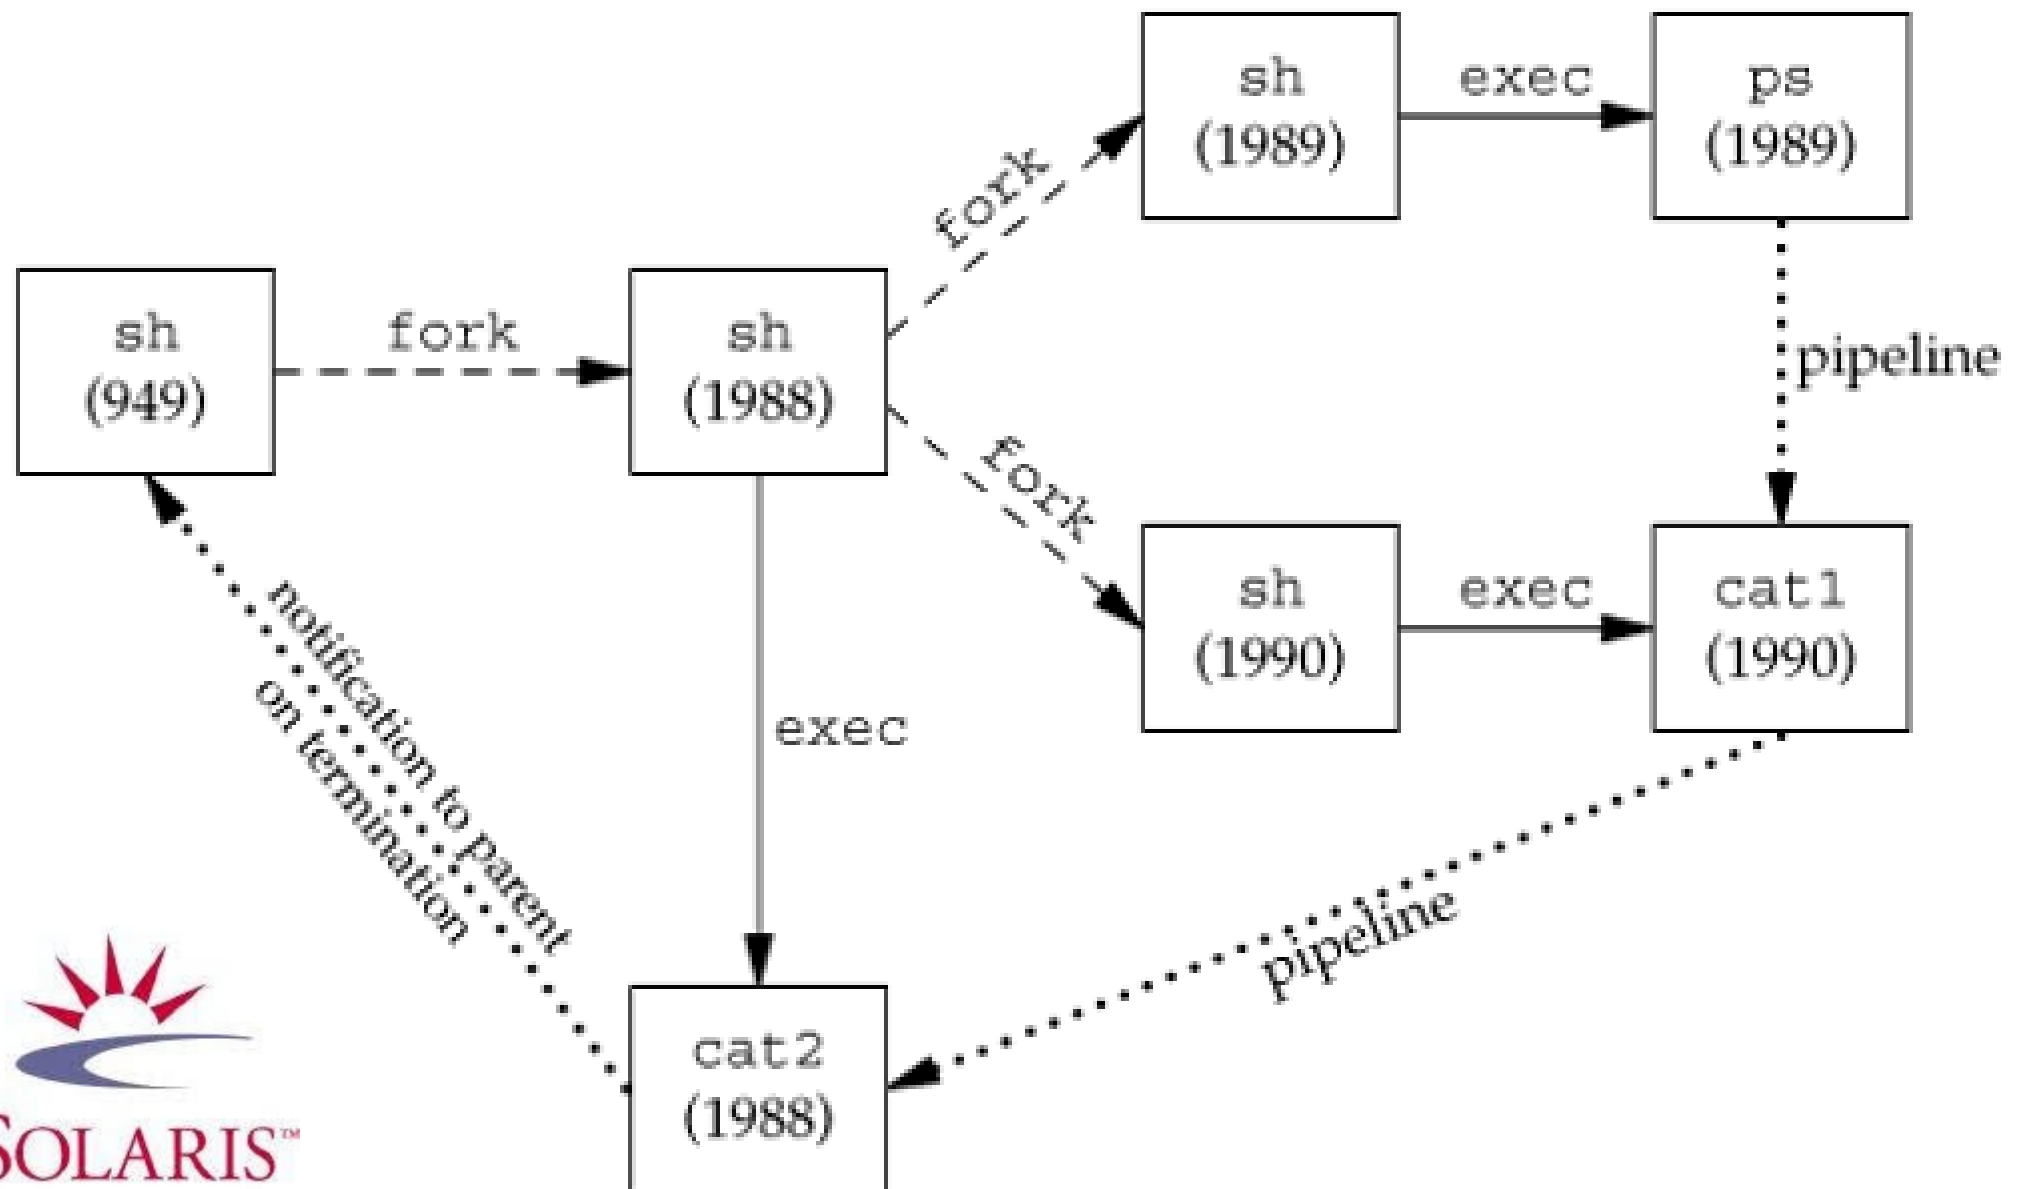
W. Richard Stevens (APUE/2E), Fig. 8.14 Differences among the six exec functions

Function	<i>pathname</i>	<i>filename</i>	Arg list	<i>argv</i> []	<i>environ</i>	<i>envp</i> []
execl	•		•		•	
execlp		•	•		•	
execle	•		•			•
execv	•			•	•	
execvp		•		•	•	
execve	•			•		•
(letter in name)		p	l	v		e

APUE/2E, Fig. 8.15. Relationship of the six exec function



APUE/2E, Fig.9.9. (Solaris 9) Processes in the pipeline **ps | cat1 | cat2** when invoked by Bourne shell



```
$ echo $SHELL  
/bin/bash
```

```
$ ls -l /bin/bash  
-rwxr-xr-x 1 root root 1113504 abr  4 13:30 /bin/bash
```

```
$ which cat  
/bin/cat
```

```
$ tee < /bin/cat cat1 > cat2
```

```
$ ls -l cat?  
-rw-rw-r-- 1 vk vk 35064 set  9 13:22 cat1  
-rw-rw-r-- 1 vk vk 35064 set  9 13:22 cat2
```

```
$ chmod a+x cat?
```

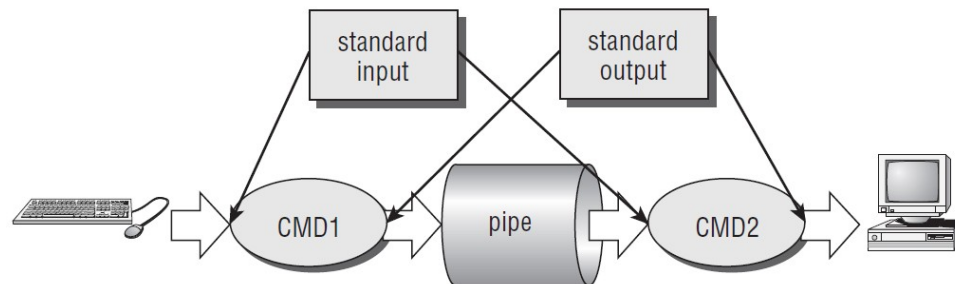
```
$ ls -l cat?  
-rwxrwxr-x 1 vk vk 35064 set  9 13:22 cat1  
-rwxrwxr-x 1 vk vk 35064 set  9 13:22 cat2
```

```
$ echo $$
```

```
17737
```

```
$ ps -o pid,ppid,pgid,sid,comm | ./cat1 | ./cat2
```

PID	PPID	PGID	SID	COMMAND
17737	16240	17737	17737	bash
18250	17737	18250	17737	ps
18251	17737	18250	17737	cat1
18252	17737	18250	17737	cat2



BLP4E


```
$ which sh
/bin/sh
```

```
$ ls -l /bin/sh
```

```
lrwxrwxrwx 1 root root 4 ago  6 20:02 /bin/sh -> dash
```

Debian Almquist Shell

```
$ which dash
/bin/dash
```

```
$ ls -l /bin/dash
```

```
-rwxr-xr-x 1 root root 121432 ene 25  2018 /bin/dash
```

```
$ dash
```

```
$ ps -o pid,ppid,pgid,sid,comm | ./cat1 | ./cat2
```

PID	PPID	PGID	SID	COMMAND
17737	19839	17737	17737	bash
18781	17737	18781	17737	dash
18814	18781	18814	17737	ps
18815	18781	18814	17737	cat1
18816	18781	18814	17737	cat2

```
$ echo $0
```

```
dash
```

```
$ ps $$
```

```
  PID TTY          STAT       TIME COMMAND
18781 pts/5        S           0:00   dash
```

```
$ cat /etc/shells
```

```
# /etc/shells: valid login shells
```

```
/bin/sh
```

```
/bin/dash
```


```
/bin/bash
```

```
/bin/rbash
```


```
$ exit
```

```
$ echo $0
```

```
bash
```



Csh
Shell with C-like syntax
★★★★☆ 4.0
1 Reseñas



Ksh
Real, AT&T version of the Korn shell
★★★★☆ 4.5
2 Reseñas



Zsh
Shell with lots of features
★★★★☆ 4.7
15 Reseñas

...



Steve Bourne

```
$ date
```

```
dom set  9 13:45:45 -05 2018
```

```
$ date > foo
```

```
$ ls -l foo
```

```
-rw-rw-r-- 1 vk vk 29 set  9 13:46 foo
```

```
$ cat foo
```

```
dom set  9 13:46:39 -05 2018
```

```
$ > foo
```

```
$ ls -l foo
```

```
-rw-rw-r-- 1 vk vk 0 set  9 13:50 foo
```

Example 4.35 (USP-CCT, Ch.4)

```
$ man cat
```

```
CAT(1)
```

User Commands

```
CAT(1)
```

NAME

cat - concatenate files and print on the standard output

SYNOPSIS

```
cat [OPTION]... [FILE]...
```

DESCRIPTION

Concatenate FILE(s) to standard output.

With no FILE, or when FILE is -, read standard input.

...

```
$ cat
```

```
Hola
```

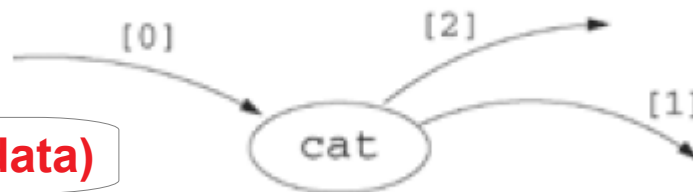
```
Hola
```

```
$
```

from standard input line by line

to standard output

Ctrl-D (end-of-data)

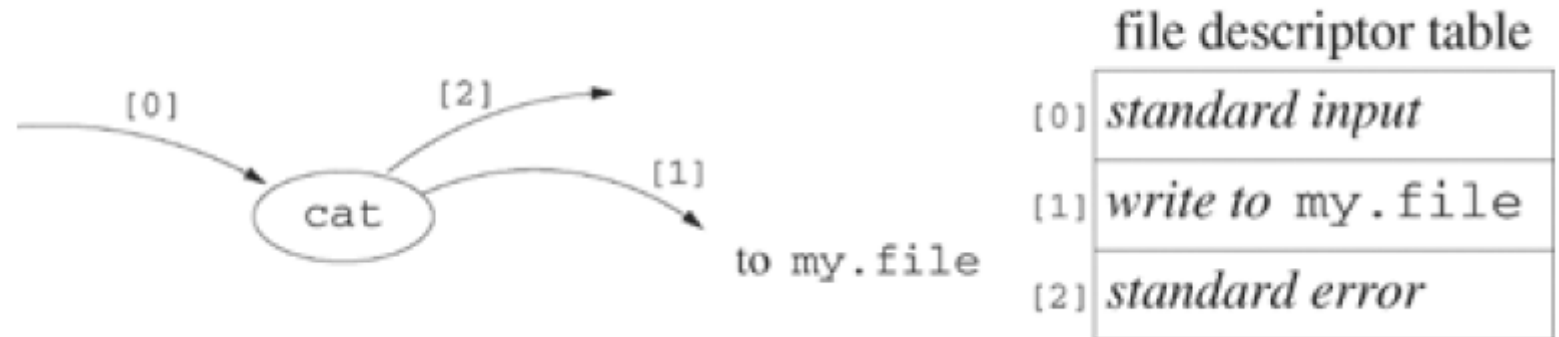


file descriptor table

[0]	<i>standard input</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>

```
$ cat > my.file
Hola
$
```

After redirection by shell and
before the cat-process execution



```
$ ls -l my.file
-rw-rw-r-- 1 vk vk 5 set  9 14:50 my.file
```


Program 4.18 (USP-CCT, Ch.4)

```
$ cat -n usp_prog4.18.c | expand
```

```
1  #include <fcntl.h>
2  #include <stdio.h>
3  #include <unistd.h>
4
5  #define CREATE_FLAGS (O_WRONLY|O_CREAT|O_APPEND)
6  #define CREATE_MODE (S_IRUSR|S_IWUSR|S_IRGRP|S_IROTH) /* rw-r--r-- */
7  int
8  main(void) {
9      int fd;
10
11      fd = open("my.file",CREATE_FLAGS,CREATE_MODE);
12      if (fd == -1) {
13          perror("Failed to create my.file");
14          return 1;
15      }
16      if (dup2(fd,STDOUT_FILENO) == -1) {
17          perror("Failed to redirect standard output");
18          return 2;
19      }
20      if (close(fd) == -1) {
21          perror("Failed to close the file");
22          return 3;
23      }
24      if (write(STDOUT_FILENO,"Ok",2) == -1) {
25          perror("Failed in writing to file");
26          return 4;
27      }
28      return 0;
29 }
```

Program 4.18 (USP-CCT, Ch.4)

```
$ gcc -o usp_prog4.18 usp_prog4.18.c
```

```
$ ./usp_prog4.18
```

```
$ ls -l my.file
```

```
-rw-r--r-- 1 vk vk 2 set  9 15:43 my.file
```

```
$ cat my.file
```

```
Ok$
```

after open

file descriptor table	
[0]	<i>standard input</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>
[3]	<i>write to my.file</i>

after dup2

file descriptor table	
[0]	<i>standard input</i>
[1]	<i>write to my.file</i>
[2]	<i>standard error</i>
[3]	<i>write to my.file</i>

after close

file descriptor table	
[0]	<i>standard input</i>
[1]	<i>write to my.file</i>
[2]	<i>standard error</i>

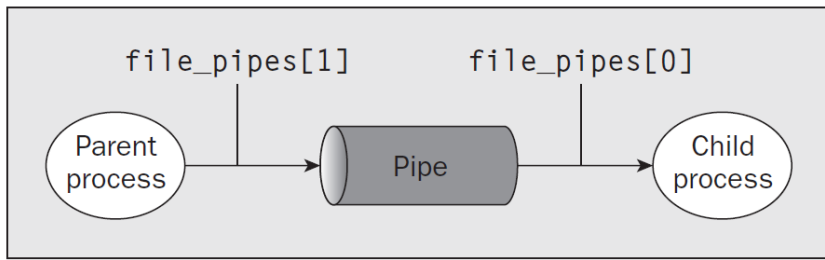


Figure 13-2

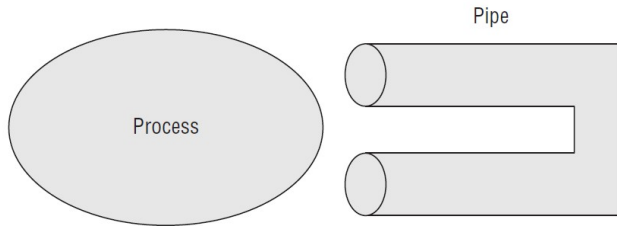


Figure 13-3

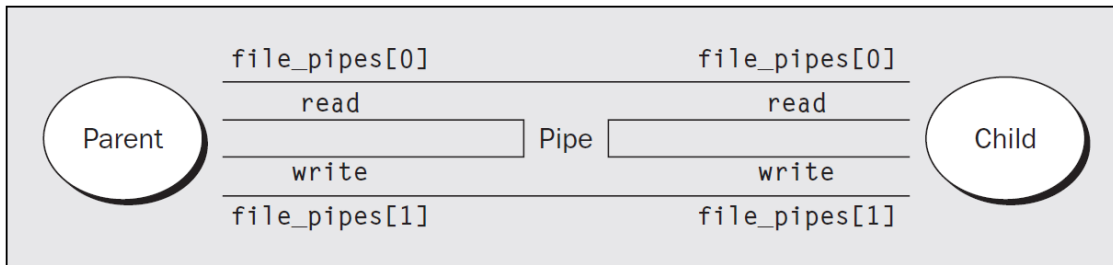


Figure 13-4

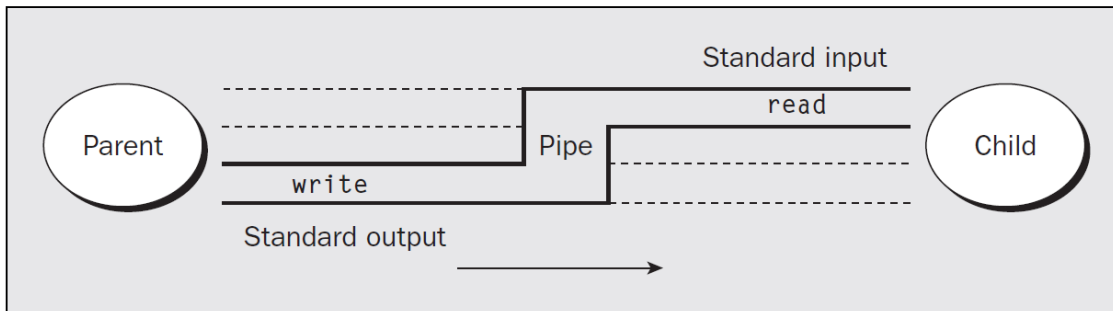


Figure 13-5

BLP4E

Preparación de un *pipe* entre dos procesos

```
#define STD_INPUT 0
#define STD_OUTPUT 1

pipeline(process1, process2)
char *process1, *process2;
{
    int fd[2];

    pipe(&fd[0]);
    if (fork() != 0) { /* El proceso padre */
        close(fd[0]);
        close(fd[1]);
        execl(process1, process1, 0);
    } else { /* El proceso hijo ejecuta esta parte */
        close(fd[1]);
        close(fd[0]);
        execl(process2, process2, 0);
    }
}
```

File descriptor	Purpose	POSIX name	<i>stdio</i> stream
0	standard input	STDIN_FILENO	<i>stdin</i>
1	standard output	STDOUT_FILENO	<i>stdout</i>
2	standard error	STDERR_FILENO	<i>stderr</i>

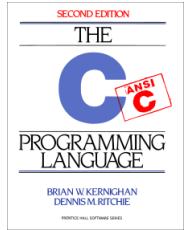
dup2(fd[1], STD_OUTPUT);

close(STD_OUTPUT); dup(fd[1]);

close(STD_INPUT); dup(fd[0]);

dup2(fd[0], STD_INPUT);

Brian W. Kernighan, Dennis M. Ritchie, El Lenguaje de Programación C, Segunda edición, Sección 7.1 del Capítulo 7



```
$ cat capital.c
#include <stdio.h>
#include <ctype.h>
#include <fcntl.h>

#define STD_IN 0

int main(int argc, char *argv[])
{
    int c, fd;

    /* si está dado el parámetro, lo consideramos como el nombre del archivo, lo abrimos para
       obtener su descriptor. Consideramos este archivo como la entrada estándar. */

    if (argc == 2) { fd = open(argv[1], O_RDONLY); close(STD_IN); dup(fd); }

    while ((c = getchar()) != EOF) putchar(toupper(c));
    return 0;
}

$ capital < capital.c      # using standard input
$ capital capital.c
```

```
$ cat mcap.c
```

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main(int argc, char **argv)
{
    int np, i, status;
    pid_t childpid;

    if (argc == 1) { printf("El formato es: %s <lista-de-archivos>\n", argv[0]);
                      exit(1); }

    np = argc - 1;          /* número de procesos a crear */
    for ( i=1; i <= np; i++ ) {
        if ( (childpid=fork()) == -1 ) { perror("fork fallo...\n"); exit(2); }
        else if ( childpid == 0 ) { /* es el hijo */
            execl("./capital", "capital", argv[i], NULL);
            perror("execl fallo...\n");
            exit(3);
        }
    }
    while (np--) wait(&status);
    return 0;
}
```

```
$ mcap capital.c capital.c capital.c
```

```
$ cat mcapsort.c
```

```
...
    int np, i, status, fd[2];
...
    pipe(fd);
    np = argc - 1;          /* número de procesos a crear */

    for (...) {
        if ( (childpid=fork()) == -1 ) { perror("fork fallo...\n"); exit(2); }
        else if ( childpid == 0 ) { /* es el hijo */
            dup2(fd[1], STDOUT_FILENO); close(fd[0]); close(fd[1]);
            execl("./capital", "capital", argv[i], NULL);
            perror("execl fallo...\n"); exit(3);
        }
    }
    while (np--) wait(&status);
    if ( (childpid=fork()) == -1 ) { perror("fork para sort fallo...\n"); exit(4); }
    else if ( childpid == 0 ) { /* el hijo para sort */
        dup2(fd[0], STDIN_FILENO); close(fd[0]); close(fd[1]);
        execl("/usr/bin/sort", "sort", NULL);
        perror("execl para sort fallo...\n");
        exit(5);
    }
    close(fd[0]); close(fd[1]); /* cerrar la tubería */
    wait(&status); /* esperar la terminación de sort */
    return 0;
}
```

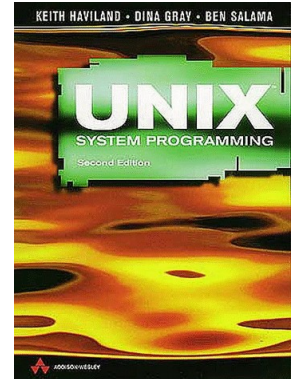

Examen 1 (2000-1), pregunta 1:

Cuando este programa termina satisfactoriamente, muestra una propiedad importante de la implementación de una tubería. ¿Cuál es esta propiedad y cuál sería el mensaje faltante en la función `alarm_action()`? (Keith Haviland, Ben Salama. *UNIX System Programming*, AW, 1997)

```
#include <signal.h>
```

```
int count;  
int alarm_action();
```

```
main()  
{  
    int p[2];  
    char c = 'x';  
  
    if (pipe(p) < 0) { perror("pipe call"); exit(1); }  
    signal(SIGALRM, alarm_action);  
    for (count = 0;;) {  
        alarm(20); /* unsigned int alarm(unsigned int seconds) */  
        write(p[1], &c, 1);  
        alarm(0); /* Alarm requests are not stacked; successive calls reset the alarm clock. */  
        /* If the argument is 0, any alarm request is canceled. */  
        if ((++count % 1024) == 0)  
            printf("%d characters in pipe\n", count);  
    }  
}  
alarm_action()  
{  
    printf(...);  
    exit(0);  
}
```



Examen 1 (2001-1), pregunta 3:

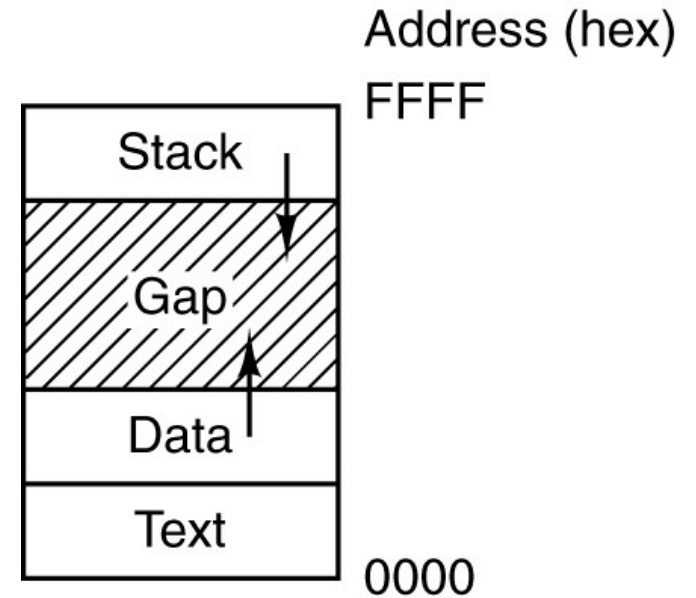
¿Cuál será el resultado de ejecución del siguiente programa?

```
$ cat pointerex.c
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main(void)
{
    int pid;
    int value=5;
    int *pval = &value;

    if ( (pid = fork() ) != 0) {
        printf("parent: value = %d, *(pval) = %d\n", value, *pval);
        value = 7;
        printf("parent: value = %d, *(pval) = %d\n", value, *pval);
    } else {
        value = 8;
        printf("child: value = %d, *(pval) = %d\n", value, *pval);
        value = 9;
        printf("child: value = %d, *(pval) = %d\n", value, *pval);
    }
}

$ gcc pointerex.c
$ a.out
...
```



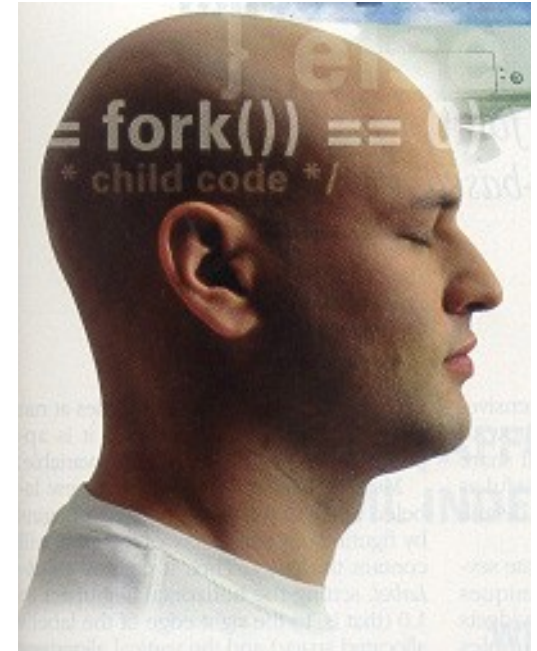
Examen 1 (2002-2), pregunta 4:

¿Cuál serán los resultados de ejecución del siguiente programa?

```
$ cat exam1-2002-2.c
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char **argv)
{
    int i, n, status;
    char str_n[8];

    if (argc == 1) exit(1);
    n = atoi(argv[1]);
    sprintf(str_n, "%d", --n);
    if (!n) execl("/bin/ps", "ps", "-l", NULL);
    if ( fork() && fork() )
        for (i=0; i<2; i++) wait(&status);
    else execl("./exam1", "exam1", str_n, NULL);
    return 0;
}
$ gcc -o exam1 exam1-2002-2.c
$ exam1
...
$ exam1 1
...
$ exam1 2
...
$ exam1 4
...
```



Examen 1 (2003-1), pregunta 4:

¿Cuál serán los resultados de ejecución del siguiente programa?

```
#include <stdio.h>
#include <unistd.h>

int main(void)
{
    int status;

    if ( (fork() && fork()) || (fork() && fork()) || fork() ) {
        while( wait(&status) != -1 );
        printf("pid = %d, ppid = %d\n", getpid(), getppid());
    }
    else {
        printf("pid = %d, ppid = %d\n", getpid(), getppid());
        sleep(1);
        execl("/bin/ps", "ps", "-l", NULL);
    }
    return 0;
}
```



M.S. Escher, *Waterfall*

Examen 2 (2003-2), pregunta 1:

¿Cuál serán los resultados de ejecución del siguiente programa?

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(void)
{
    int i=0, status;

    fprintf(stderr, "pid = %ld, ppid = %ld.\n",
              (long) getpid(), (long) getppid());
    while ( (!fork() || !fork()) && i<2 ) {
        fprintf(stderr, "pid = %ld, ppid = %ld: %d\n",
              (long) getpid(), (long) getppid(), i);
        if ( i % 2 ) break;
        i++;
    }
    sleep(1);
    while(waitpid(&status) != -1);
    fprintf(stderr, "pid = %ld, ppid = %ld: terminated.\n",
              (long) getpid(), (long) getppid());
    return 0;
}
```



Examen 1 (2004-1), pregunta 1:

¿Cuál serán los resultados de ejecución del siguiente programa?

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <signal.h>

int main(void)
{
    pid_t great_grandfather, ppid, bill, bill_vol2;
    int status, pipefd[2];

    pipe(pipefd);
    dup2(pipefd[0],STDIN_FILENO); dup2(pipefd[1],STDOUT_FILENO);
    close(pipefd[0]); close(pipefd[1]);
    fprintf(stderr, "pid = %ld, ppid = %ld.\n",
        (long)(great_grandfather=getpid()), (long)getppid());
    if ( (!fork() || fork()) && (fork() == !(bill=fork())) ) {
        fprintf(stderr, "pid = %ld, ppid = %ld: true and bill = %ld\n",
            (long)getpid(), (long)(ppid=getppid()), (long)bill);
        sleep(1);
        if (ppid != great_grandfather) {
            write(1, &bill, sizeof(ppid));
            dup2(STDERR_FILENO,STDOUT_FILENO);
            execl("/bin/ps", "ps", "-l", NULL);
        } else {
            ...
        }
    }
}
```



```

...
} else {
    bill_vol2 = bill;
    read(0, &bill, sizeof(ppid));
    fprintf(stderr, "pid = %ld, ppid = %ld: father of %ld did ps\n",
        (long) getpid(), (long) getppid(), bill);
    kill(bill, SIGTERM);
    fprintf(stderr, "pid = %ld: terminated.\n", bill);
    kill(bill_vol2, SIGTERM);
    fprintf(stderr, "pid = %ld: terminated.\n", bill_vol2);
    fprintf(stderr, "pid = %ld, ppid = %ld: finished.\n",
        (long) getpid(), (long) getppid());
    exit(0);
}
}
close(0); close(1);
sleep(3);
while (waitpid(-1, &status, 0) != -1);
fprintf(stderr, "pid = %ld, ppid = %ld: finished.\n",
    (long) getpid(), (long) getppid());
return 0;
}

```

