

Programming Python by Mark Lutz

Chapter 5: Parallel System Tools

Forking Processes

"forks child processes until you type 'q'"

```
import os
```

```
def child():  
    print('Hello from child', os.getpid())  
    os._exit(0)  # else goes back to parent loop
```

```
def parent():  
    while True:  
        newpid = os.fork()  
        if newpid == 0:  
            child()  
        else:  
            print('Hello from parent', os.getpid(), newpid)  
            if input() == 'q': break
```

```
parent()
```

PP4E/System/Processes/fork1.py

```
$ python3 fork1.py
Hello from parent 657 658
Hello from child 658

Hello from parent 657 659
Hello from child 659

Hello from parent 657 660
Hello from child 660

Hello from parent 657 661
Hello from child 661

Hello from parent 657 662
Hello from child 662

Hello from parent 657 663
Hello from child 663

q
$
```

PP4E/System/Processes/fork-count.py

```
"""
fork basics: start 5 copies of this program running in parallel with
the original; each copy counts up to 5 on the same stdout stream--forks
copy process memory, including file descriptors; fork doesn't currently
work on Windows without Cygwin: use os.spawnv or multiprocessing on
Windows instead; spawnv is roughly like a fork+exec combination;
"""

import os, time

def counter(count):
    for i in range(count):
        time.sleep(1)
        print('[%s] => %s' % (os.getpid(), i))

for i in range(5):
    pid = os.fork()
    if pid != 0:
        print('Process %d spawned' % pid)
    else:
        counter(5)
process
    os._exit(0)

print('Main process exiting.')
```

```
$ python3 fork-count.py
```

```
Process 837 spawned
```

```
Process 838 spawned
```

```
Process 839 spawned
```

```
Process 840 spawned
```

```
Process 841 spawned
```

```
Main process exiting.
```

```
$ [837] => 0
```

```
[838] => 0
```

```
[839] => 0
```

```
[840] => 0
```

```
[841] => 0
```

```
...
```

```
[837] => 4
```

```
[838] => 4
```

```
[839] => 4
```

```
[840] => 4
```

```
[841] => 4
```

PP4E/System/Processes/fork-exec.py

```
"starts programs until you type 'q'"
```

```
import os
```

```
parm = 0
```

```
while True:
```

```
    parm += 1
```

```
    pid = os.fork()
```

```
    if pid == 0:
```

```
        os.execvp('python3', 'python3', 'child.py', str(parm)) # copy process  
        assert False, 'error starting program' # overlay program  
                                                # shouldn't return
```

```
else:
```

```
    print('Child is', pid)
```

```
    if input() == 'q': break
```

PP4E/System/Processes/child.py

```
import os, sys

print('Hello from child', os.getpid(), sys.argv[1])
```

PP4E/System/Processes/fork-exec.py

```
$ python3 fork-exec.py
Child is 897
('Hello from child', 897, '1')
q
```

```
$ python3 fork-exec.py
Child is 929
Hello from child 929 1
```

```
Child is 930
Hello from child 930 2
```

```
Child is 931
Hello from child 931 3
```

```
Child is 932
Hello from child 932 4
```

```
Child is 933
Hello from child 933 5
```

```
q
```


PP4E/System/Processes/pipe1.py

```
import os, time

def child(pipeout):
    zzz = 0
    while True:
        time.sleep(zzz)
        msg = ('Spam %03d' % zzz).encode()
        os.write(pipeout, msg)
        zzz = (zzz+1) % 5
        # make parent wait
        # pipes are binary bytes
        # send to parent
        # goto 0 after 4

def parent():
    pipein, pipeout = os.pipe()
    if os.fork() == 0:
        child(pipeout)
    else:
        # make 2-ended pipe
        # copy this process
        # in copy, run child
        # in parent, listen to pipe
        while True:
            line = os.read(pipein, 32)
            print('Parent %d got [%s] at %s' % (os.getpid(), line, time.time()))
            # blocks until data sent

parent()
```

```
$ python3 pipel.py
Parent 4391 got [b'Spam 000'] at 1332781761.500339
Parent 4391 got [b'Spam 001'] at 1332781762.501251
Parent 4391 got [b'Spam 002'] at 1332781764.502363
Parent 4391 got [b'Spam 003'] at 1332781767.505495
Parent 4391 got [b'Spam 004Spam 000'] at 1332781771.508707
Parent 4391 got [b'Spam 001'] at 1332781772.509974
Parent 4391 got [b'Spam 002'] at 1332781774.512041
Parent 4391 got [b'Spam 003'] at 1332781777.512628
Parent 4391 got [b'Spam 004Spam 000'] at 1332781781.516459
Parent 4391 got [b'Spam 001'] at 1332781782.516878
Parent 4391 got [b'Spam 002'] at 1332781784.51905
...
```

PP4E/System/Processes/pipe2.py

```
# same as pipe1.py, but wrap pipe input in stdio file object
# to read by line, and close unused pipe fds in both processes

import os, time

def child(pipeout):
    zzz = 0
    while True:
        time.sleep(zzz)
        msg = ('Spam %03d\n' % zzz).encode()
        os.write(pipeout, msg)
        zzz = (zzz+1) % 5
        # make parent wait
        # pipes are binary in 3.X
        # send to parent
        # roll to 0 at 5

def parent():
    pipein, pipeout = os.pipe()
    if os.fork() == 0:
        os.close(pipein)
        child(pipeout)
    else:
        os.close(pipeout)
        pipein = os.fdopen(pipein)
        while True:
            line = pipein.readline()[:-1]
            print('Parent %d got [%s] at %s' % (os.getpid(), line, time.time()))
            # in parent, listen to pipe
            # close output side here
            # make text mode input file object
            # blocks until data sent

parent()
```

```
$ python3 pipe2.py
Parent 4414 got [Spam 000] at 1332782205.789812
Parent 4414 got [Spam 001] at 1332782206.791198
Parent 4414 got [Spam 002] at 1332782208.792644
Parent 4414 got [Spam 003] at 1332782211.795856
Parent 4414 got [Spam 004] at 1332782215.796212
Parent 4414 got [Spam 000] at 1332782215.796337
Parent 4414 got [Spam 001] at 1332782216.797505
Parent 4414 got [Spam 002] at 1332782218.798857
Parent 4414 got [Spam 003] at 1332782221.800243
Parent 4414 got [Spam 004] at 1332782225.804157
Parent 4414 got [Spam 000] at 1332782225.804266
Parent 4414 got [Spam 001] at 1332782226.805408
Parent 4414 got [Spam 002] at 1332782228.807614
...
```

PP4E/System/Processes/pipes.py

```
"""
```

```
spawn a child process/program, connect my stdin/stdout to child process's
stdout/stdin--my reads and writes map to output and input streams of the
spawned program; much like tying together streams with subprocess module;
```

```
"""
```

```
import os, sys
```

```
def spawn(prog, *args):
    stdinFd = sys.stdin.fileno()
    stdoutFd = sys.stdout.fileno()

    parentStdin, childStdout = os.pipe()
    childStdin, parentStdout = os.pipe()
    pid = os.fork()
    if pid:
        os.close(childStdout)
        os.close(childStdin)
        os.dup2(parentStdin, stdinFd)
        os.dup2(parentStdout, stdoutFd)
    else:
        os.close(parentStdin)
        os.close(parentStdout)
        os.dup2(childStdin, stdinFd)
        os.dup2(childStdout, stdoutFd)
        args = (prog,) + args
        os.execvp(prog, args)
        assert False, 'execvp failed!'
```

```
# pass progname, cmdline args
# get descriptors for streams
# normally stdin=0, stdout=1

# make two IPC pipe channels
# pipe returns (inputfd, outputfd)
# make a copy of this process

# in parent process after fork:
# close child ends in parent
# my sys.stdin copy = pipe1[0]
# my sys.stdout copy = pipe2[1]

# in child process after fork:
# close parent ends in child
# my sys.stdin copy = pipe2[0]
# my sys.stdout copy = pipe1[1]

# new program in this process
# os.exec call never returns here
```

```
if __name__ == '__main__':
    mypid = os.getpid()
    spawn('python3', 'pipes-testchild.py', 'spam')      # fork child program

    print('Hello 1 from parent', mypid)                  # to child's stdin
    sys.stdout.flush()                                   # subvert stdio buffering
    reply = input()                                      # from child's stdout
    sys.stderr.write('Parent got: "%s"\n' % reply)        # stderr not tied to pipe!

    print('Hello 2 from parent', mypid)
    sys.stdout.flush()
    reply = sys.stdin.readline()
    sys.stderr.write('Parent got: "%s"\n' % reply[:-1])
```

PP4E/System/Processes/pipes-testchild.py

```
import os, time, sys

mypid      = os.getpid()
parentpid = os.getppid()
sys.stderr.write('Child %d of %d got arg: "%s"\n' %
                 (mypid, parentpid, sys.argv[1]))

for i in range(2):
    time.sleep(3)          # make parent process wait by sleeping here
    recv = input()         # stdin tied to pipe: comes from parent's stdout
    time.sleep(3)
    send = 'Child %d got: [%s]' % (mypid, recv)
    print(send)            # stdout tied to pipe: goes to parent's stdin
    sys.stdout.flush()     # make sure it's sent now or else process blocks
```

```
$ python3 pipes.py
Child 4481 of 4480 got arg: "spam"
Parent got: "Child 4481 got: [Hello 1 from parent 4480]"
Parent got: "Child 4481 got: [Hello 2 from parent 4480]"
$
```