

open + ... + mount + ...

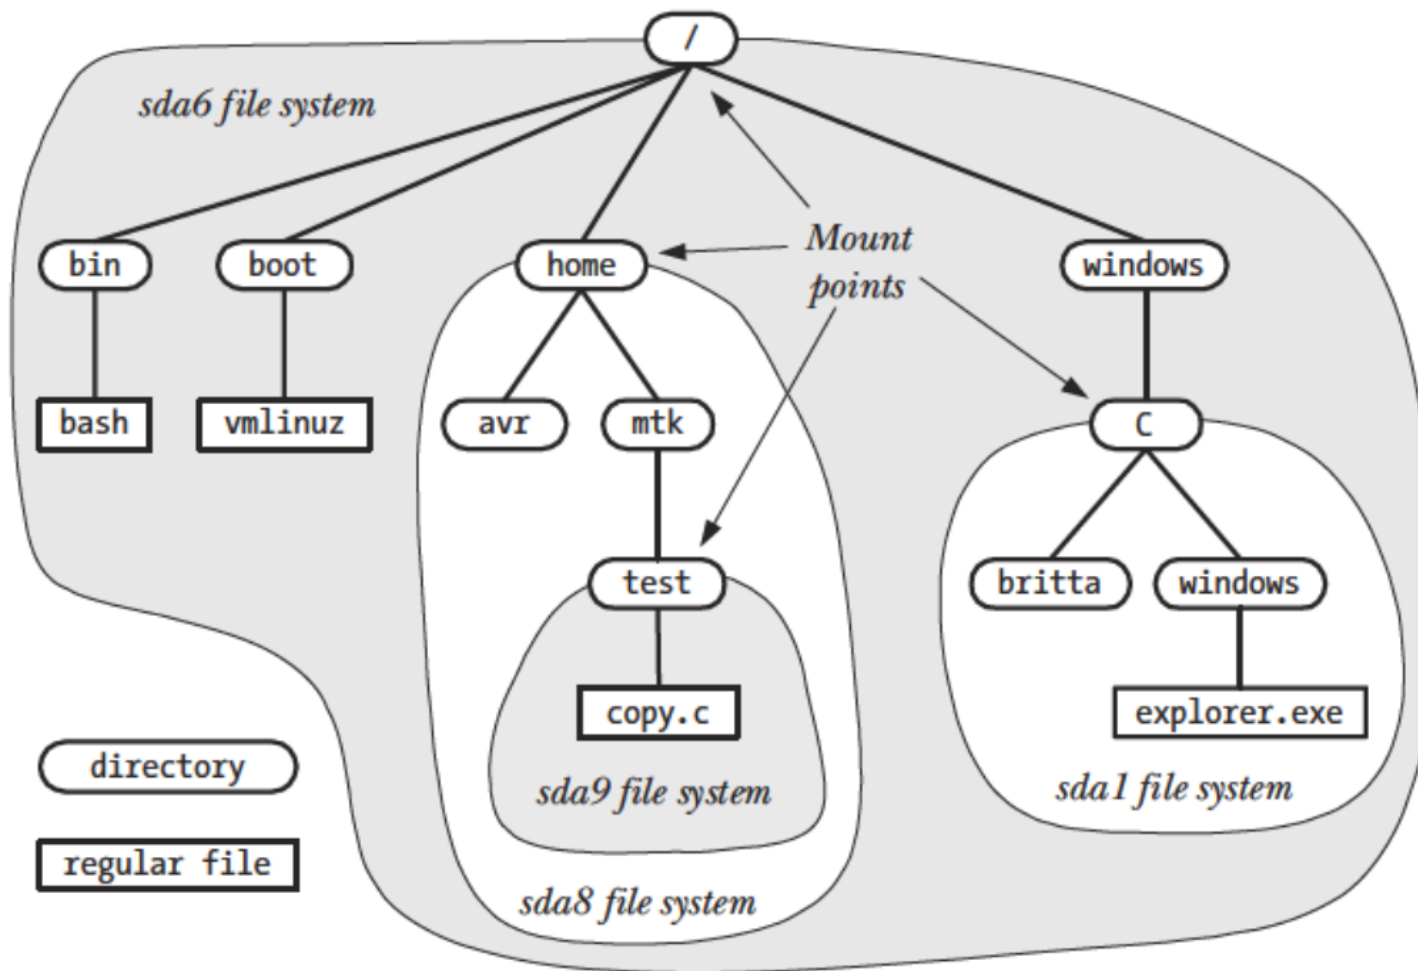
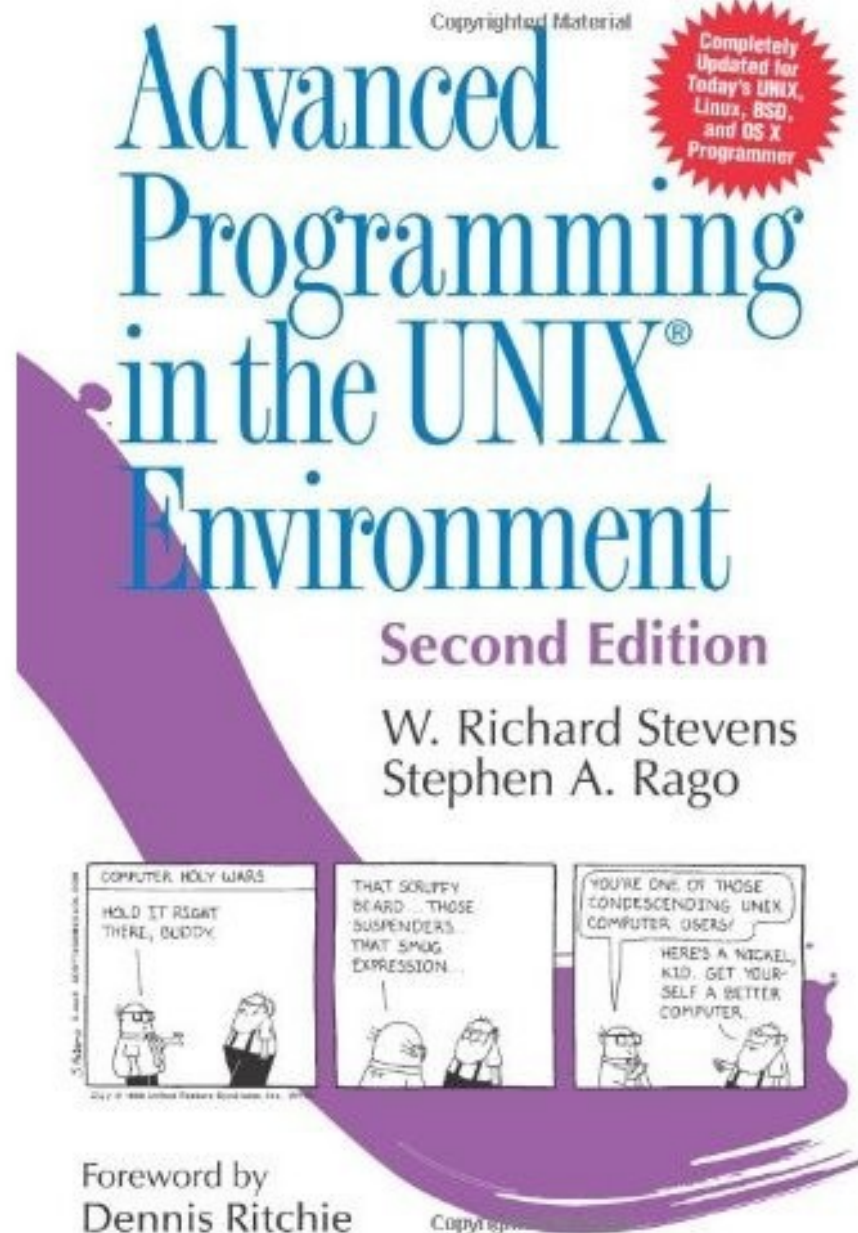


Figure 14-4: Example directory hierarchy showing file-system mount points



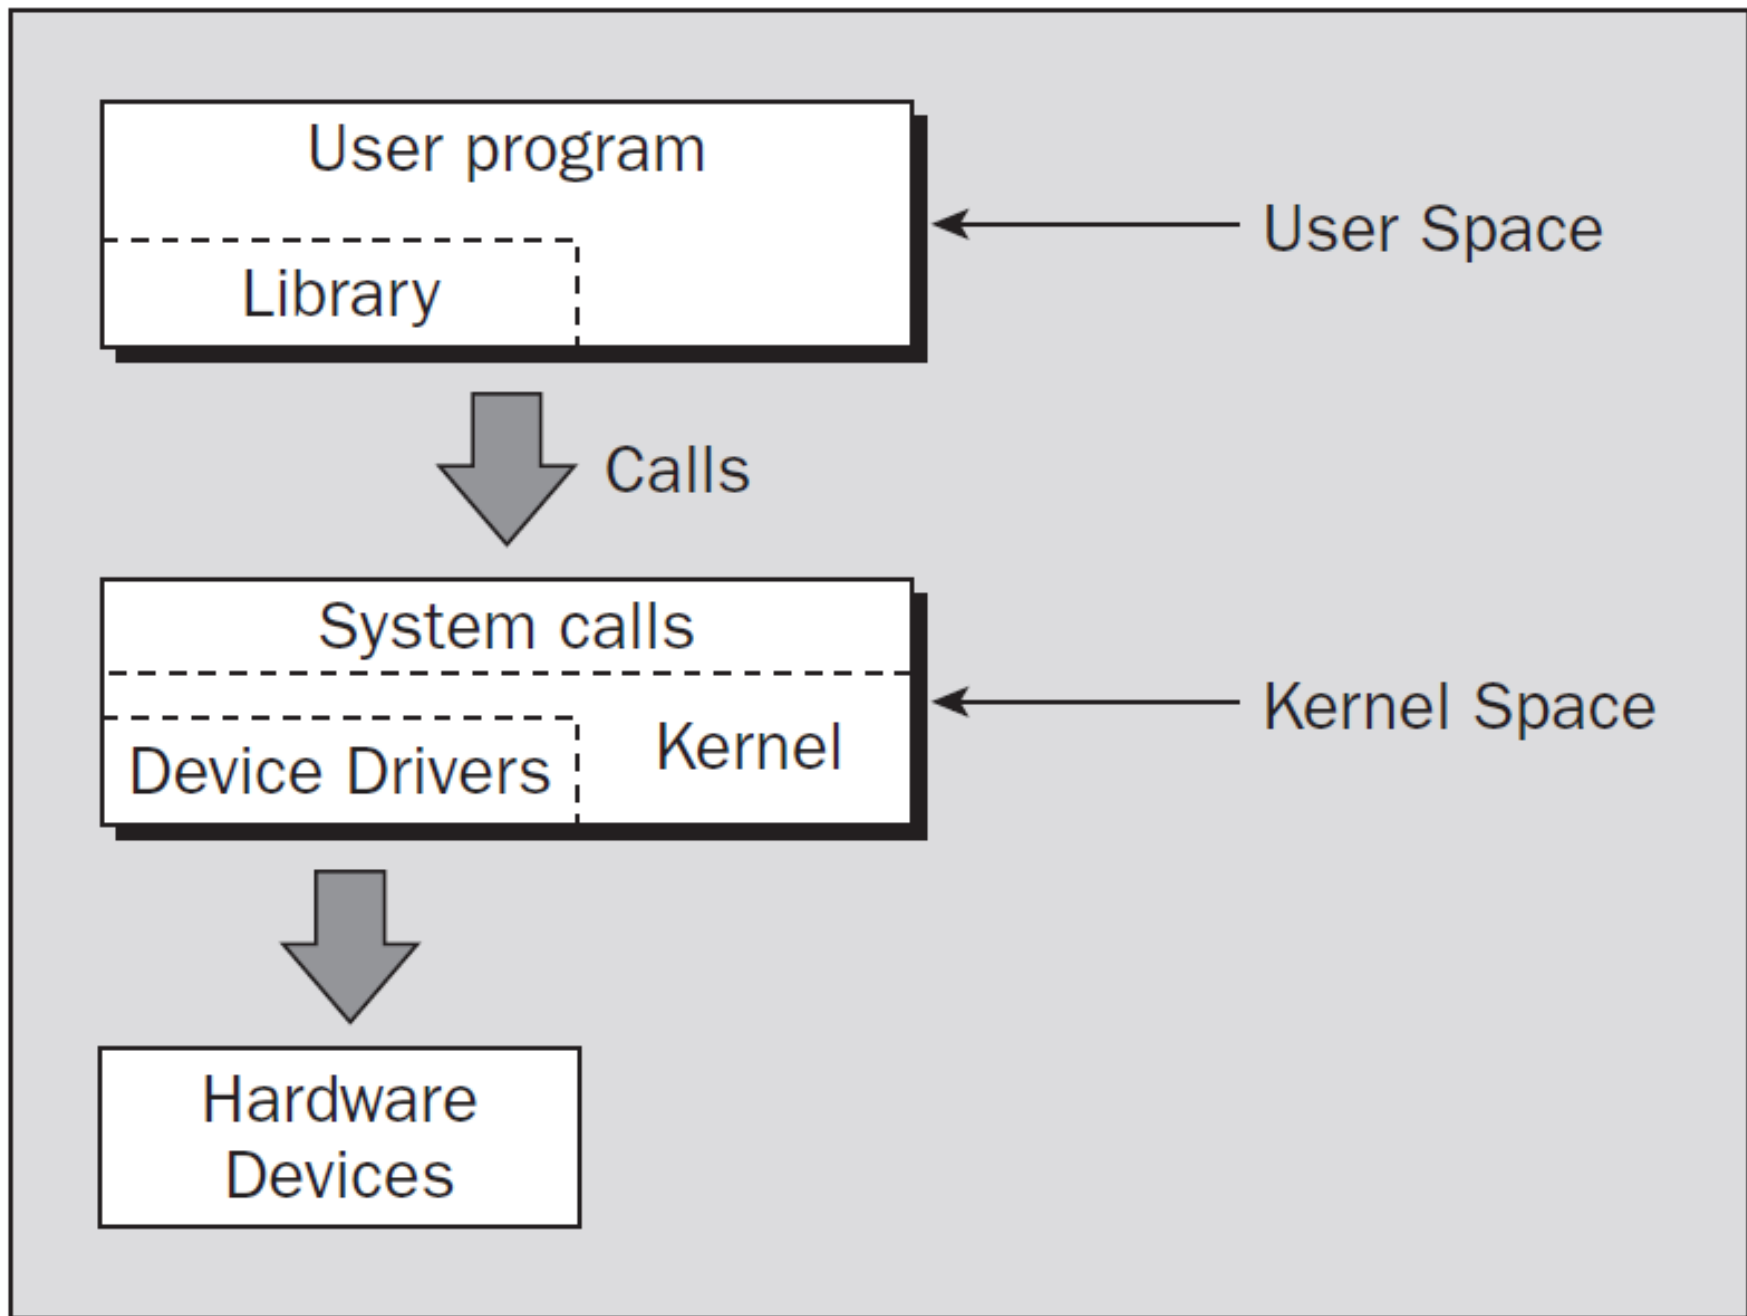
ADDISON-WESLEY PROFESSIONAL COMPUTING SERIES

THE **LINUX** PROGRAMMING INTERFACE


A Linux and UNIX® System Programming Handbook

MICHAEL KERRISK





A file is opened or created by calling the **open** function:



```
#include <fcntl.h>
```

```
int open(const char *pathname, int oflag, ... /* mode_t mode */ );
```

Returns: file descriptor if OK, -1 on error

One and only one of these three constants: **O_RDONLY**, **O_WRONLY**, **O_RDWR**.

Optional: **O_CREAT**, **O_EXCL**, **O_TRUNC**, **O_NOCTTY**, **O_NONBLOCK**, ...

An open file is closed by calling the **close** function:

```
#include <unistd.h>
```

```
int close(int fildes);
```

Returns: 0 if OK, -1 on error

An open file's offset can be set explicitly by calling **lseek**:

```
#include <unistd.h>
```

```
off_t lseek(int filedes, off_t offset, int whence);
```

Returns: new file offset if OK, -1 on error

The value of the *whence* argument: **SEEK_SET**, **SEEK_CUR**, **SEEK_END**.

Data is read from an open file with the **read** function:

```
#include <unistd.h>
```

```
ssize_t read(int filedes, void *buf, size_t nbytes);
```

Returns: number of bytes read, 0 if end of file, -1 on error

read can block the caller forever if data isn't present with certain file types (pipes, terminal devices, and network devices).

If we **read** from a pipe whose **write end** has been closed, **read** returns 0 to indicate an **end of file** after all the data has been read. We should say that this end of file is not generated until there are no more writers for the pipe. It's possible to duplicate a pipe descriptor so that multiple processes have the pipe open for writing.

Data is written to an open file with the **write** function:

```
#include <unistd.h>
```

```
ssize_t write(int filedes, const void *buf, size_t nbytes);
```

Returns: number of bytes written if OK, -1 on error

write can block the caller forever if the data can't be accepted immediately by certain file types (no room in the pipe, network flow control, etc.).

If we **write** to a pipe whose **read end** has been closed, the signal **SIGPIPE** is generated. If we either ignore the signal or catch it and return from the signal handler, **write** returns -1 with **errno** set to **EPIPE**.

File's status information can be obtained with the **stat** function:

```
#include <sys/stat.h>
```

```
int stat(const char *restrict pathname, struct stat *restrict buf);  
int fstat(int filedes, struct stat buf);  
int lstat(const char *restrict pathname, struct stat *restrict buf);
```

All three return: 0 if OK, -1 on error

```
struct stat {
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t     st_mode;     /* protection */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t  st_blksize;  /* blocksize for file system I/O */
    blkcnt_t   st_blocks;   /* number of 512B blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

Directories are created with the **mkdir** function:

```
#include <sys/stat.h>
```

```
int mkdir(const char *pathname, mode_t mode);
```

Returns: 0 if OK, -1 on error

An empty directory is deleted with the **rmdir** function:

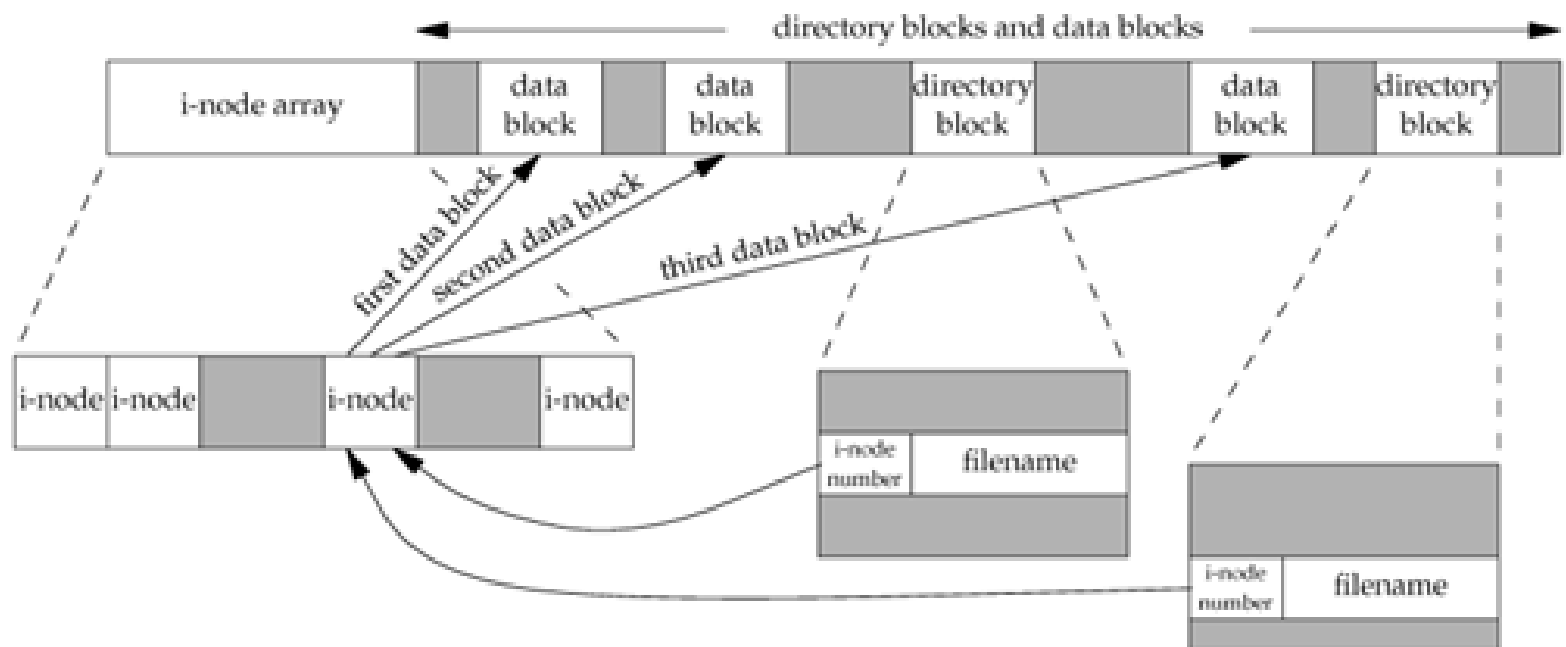
```
#include <unistd.h>
```

```
int rmdir(const char *pathname);
```

Returns: 0 if OK, -1 on error

Hard links:

Figure 4.14. Cylinder group's i-nodes and data blocks in more detail



Any file can have multiple directory entries pointing to its i-node. The way we create a link to an existing file is with the **link** function:

```
#include <unistd.h>
```

```
int link(const char *existingpath, const char *newpath);
```

Returns: 0 if OK, -1 on error

To remove an existing directory entry, we call the **unlink** function:

```
#include <unistd.h>
```

```
int unlink(const char *pathname);
```

Returns: 0 if OK, -1 on error

A **symbolic link** is an indirect pointer to a file.

A symbolic link is created with the **symlink** function:

```
#include <unistd.h>
```

```
int symlink(const char *actualpath, const char *sympath);
```

Returns: 0 if OK, -1 on error

A new directory entry, **sympath**, is created that points to **actualpath**. It is not required that **actualpath** exist when the symbolic link is created.

Because the **open** function follows a symbolic link, we need a way to open the link itself and read the name in the link. The **readlink** function does this.

```
#include <unistd.h>
```

```
ssize_t readlink(const char* restrict pathname, char *restrict buf,  
                size_t bufsize);
```

Returns: number of bytes read if OK, -1 on error

Examen 2 (2011-1), parte de la pregunta 3:

```
$ mkdir d0 d1 d2
$ echo File a > d0/a
$ ln -s a d0/b
$ ln -s ../d1/c d0/c
$ ln -s ../d0/d d1/c
$ ln d0/a d0/d
$ ln -s ../d2/f d0/e
$ echo File e > d0/e
```

Indique los contenidos de los 3 directorios creados y los contenidos de sus archivos.

```
$ mkdir d0 d1 d2
```

```
$ ls -l
```

```
total 12
```

```
drwxrwxr-x 2 vk vk 4096 set 9 16:18 d0
```

```
drwxrwxr-x 2 vk vk 4096 set 9 16:18 d1
```

```
drwxrwxr-x 2 vk vk 4096 set 9 16:18 d2
```

```
$ echo File a > d0/a
```

```
$ ls -l d0/
```

```
total 4
```

```
-rw-rw-r-- 1 vk vk 7 set 9 16:18 a
```

```
$ cat d0/a
```

```
File a
```



```
$ ln -s a d0/b
```

```
$ ls -l d0/
```

```
total 4
```

```
-rw-rw-r-- 1 vk vk 7 set  9 16:18 a
```

```
lrwxrwxrwx 1 vk vk 1 set  9 16:20 b -> a
```

```
$ ls -l d0/b
```

```
lrwxrwxrwx 1 vk vk 1 set  9 16:20 d0/b -> a
```

```
$ cat d0/b
```

```
File a
```

```
$ ln -s ../d1/c d0/c
```

```
$ ls -l d0/
```

```
total 4
```

```
-rw-rw-r-- 1 vk vk 7 set  9 16:18 a
```

```
lrwxrwxrwx 1 vk vk 1 set  9 16:20 b -> a
```

```
lrwxrwxrwx 1 vk vk 7 set  9 16:21 c -> ../d1/c
```

dead link

```
$ ls -l d0/c
```

```
lrwxrwxrwx 1 vk vk 7 set  9 16:21 d0/c -> ../d1/c
```

```
$ cat d0/c
```

```
cat: d0/c: No existe el archivo o el directorio
```

```
$ ln -s ../d0/d d1/c
```

```
$ ls -l d1/
```

```
total 0
```

```
lrwxrwxrwx 1 vk vk 7 set 9 16:31 c -> ../d0/d
```

```
$ ln d0/a d0/d
```

```
$ ls -li d0/
```

```
total 8
```

```
6947875 -rw-rw-r-- 2 vk vk 7 set 9 16:18 a
```

```
6947876 lrwxrwxrwx 1 vk vk 1 set 9 16:20 b -> a
```

```
6947877 lrwxrwxrwx 1 vk vk 7 set 9 16:21 c -> ../d1/c
```

```
6947875 -rw-rw-r-- 2 vk vk 7 set 9 16:18 d
```



```
$ ln -s ../d2/f d0/e
```

```
$ ls -l d0/
```

```
total 8
```

```
-rw-rw-r-- 2 vk vk 7 set 9 16:18 a  
lrwxrwxrwx 1 vk vk 1 set 9 16:20 b -> a  
lrwxrwxrwx 1 vk vk 7 set 9 16:21 c -> ../d1/c  
-rw-rw-r-- 2 vk vk 7 set 9 16:18 d  
lrwxrwxrwx 1 vk vk 7 set 9 16:33 e -> ../d2/f
```

```
$ echo File e > d0/e
```

```
$ ls -l d0/
```

```
total 8
```

```
-rw-rw-r-- 2 vk vk 7 set 9 16:18 a  
lrwxrwxrwx 1 vk vk 1 set 9 16:20 b -> a  
lrwxrwxrwx 1 vk vk 7 set 9 16:21 c -> ../d1/c  
-rw-rw-r-- 2 vk vk 7 set 9 16:18 d  
lrwxrwxrwx 1 vk vk 7 set 9 16:33 e -> ../d2/f
```

```
$ ls -l d1/
```

```
total 0
```

```
lrwxrwxrwx 1 vk vk 7 set  9 16:31 c -> ../d0/d
```

```
$ ls -l d2/
```

```
total 4
```

```
-rw-rw-r-- 1 vk vk 7 set  9 16:34 f
```

```
$ cat d0/{a,b,c,d,e}
```

```
File a
```

```
File a
```

```
File a
```

```
File a
```

```
File e
```

```
$ cat d1/c
```

```
File a
```

```
$ cat d2/f
```

```
File e
```

```

$ mount
/dev/sda6 on / type ext4 (rw)
proc on /proc type proc (rw)
sysfs on /sys type sysfs (rw)
devpts on /dev/pts type devpts (rw,mode=0620,gid=5)
/dev/sda8 on /home type ext3 (rw,acl,user_xattr)
/dev/sda1 on /windows/C type vfat (rw,noexec,nosuid,nodev)
/dev/sda9 on /home/mtk/test type reiserfs (rw)

```

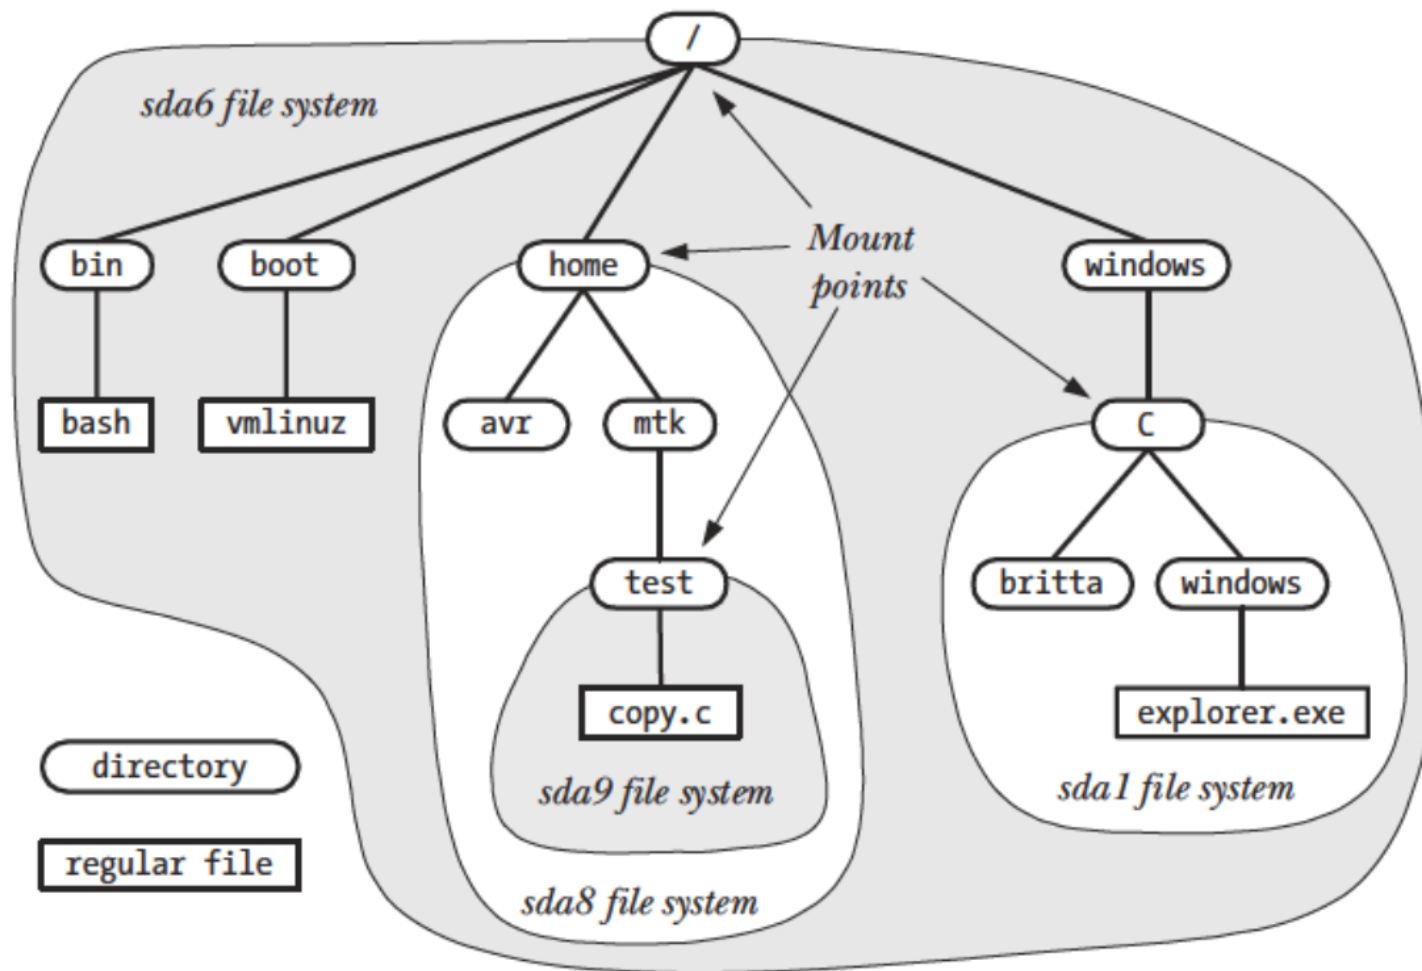


Figure 14-4: Example directory hierarchy showing file-system mount points

The **mount()** system call mounts the file system contained on the device specified by **source** under the directory (the *mount point*) specified by **target**:

```
#include <sys/mount.h>
```

```
int mount(const char *source, const char *target, const char *fstype,  
unsigned long mountflags, const void *data);
```

Returns: 0 if OK, -1 on error

The **umount()** system call unmounts a mounted file system:

```
#include <sys/mount.h>
```

```
int umount(const char *target);
```

Returns: 0 if OK, -1 on error

Every process has a current working directory. We can change the current working directory of the calling process by calling the `chdir` or `fchdir` functions:

```
#include <unistd.h>
```

```
int chdir(const char *pathname);  
int fchdir(int filedes);
```

Both return: 0 if OK, -1 on error

These two functions allow us to change the file access permissions for an existing file:

```
#include <sys/stat.h>
```

```
int chmod(const char *pathname, mode_t mode);  
int fchmod(int fildes, mode_t mode);
```

Both return: 0 if OK, -1 on error

The **kill** function sends a signal to a process or a group of processes:

```
#include <signal.h>
```

```
int kill(pid_t pid, int signo);
```

Returns: 0 if OK, -1 on error

\$ man 7 signal



...

Signal	Value	Action	Comment
SIGHUP	1	Term	Hangup detected on controlling terminal or death of controlling process
SIGINT	2	Term	Interrupt from keyboard
SIGQUIT	3	Core	Quit from keyboard
SIGILL	4	Core	Illegal Instruction
SIGABRT	6	Core	Abort signal from abort(3)
SIGFPE	8	Core	Floating point exception
SIGKILL	9	Term	Kill signal
SIGSEGV	11	Core	Invalid memory reference
SIGPIPE	13	Term	Broken pipe: write to pipe with no readers
SIGALRM	14	Term	Timer signal from alarm(2)
SIGTERM	15	Term	Termination signal
SIGUSR1	30,10,16	Term	User-defined signal 1
SIGUSR2	31,12,17	Term	User-defined signal 2
SIGCHLD	20,17,18	Ign	Child stopped or terminated
SIGCONT	19,18,25	Cont	Continue if stopped
SIGSTOP	17,19,23	Stop	Stop process
SIGTSTP	18,20,24	Stop	Stop typed at tty
SIGTTIN	21,21,26	Stop	tty input for background process
SIGTTOU	22,22,27	Stop	tty output for background process

The signals SIGKILL and SIGSTOP cannot be caught, blocked, or ignored.

```
$ kill -l
```

1) SIGHUP	2) SIGINT	3) SIGQUIT	4) SIGILL	5) SIGTRAP
6) SIGABRT	7) SIGBUS	8) SIGFPE	9) SIGKILL	10) SIGUSR1
11) SIGSEGV	12) SIGUSR2	13) SIGPIPE	14) SIGALRM	15) SIGTERM
16) SIGSTKFLT	17) SIGCHLD	18) SIGCONT	19) SIGSTOP	20) SIGTSTP
21) SIGTTIN	22) SIGTTOU	23) SIGURG	24) SIGXCPU	25) SIGXFSZ
26) SIGVTALRM	27) SIGPROF	28) SIGWINCH	29) SIGIO	30) SIGPWR
31) SIGSYS	34) SIGRTMIN	35) SIGRTMIN+1	36) SIGRTMIN+2	37) SIGRTMIN+3
38) SIGRTMIN+4	39) SIGRTMIN+5	40) SIGRTMIN+6	41) SIGRTMIN+7	42) SIGRTMIN+8
43) SIGRTMIN+9	44) SIGRTMIN+10	45) SIGRTMIN+11	46) SIGRTMIN+12	47) SIGRTMIN+13
48) SIGRTMIN+14	49) SIGRTMIN+15	50) SIGRTMAX-14	51) SIGRTMAX-13	52) SIGRTMAX-12
53) SIGRTMAX-11	54) SIGRTMAX-10	55) SIGRTMAX-9	56) SIGRTMAX-8	57) SIGRTMAX-7
58) SIGRTMAX-6	59) SIGRTMAX-5	60) SIGRTMAX-4	61) SIGRTMAX-3	62) SIGRTMAX-2
63) SIGRTMAX-1	64) SIGRTMAX			



Daniel Stori (<http://turnoff.us/geek/>)



Daniel Stori (<http://turnoff.us/geek/>)



Daniel Stori (<http://turnoff.us/geek/>)

On the other hand, when we brutally kill a process with a SIGKILL, we prevent them from finishing their job and say goodbye...



...and this is so SAD!



Daniel Stori (<http://turnoff.us/geek/>)



Dad, where
are we
going?

So please, DON'T use
SIGKILL. Give the kids
the chance to leave the
kernel in peace.

Be nice.


Dad, where
are you?



Daniel Stori {turnoff.us}

Daniel Stori (<http://turnoff.us/geek/>)

\$ Adopt a good cause, DON'T SIGKILL



Daniel Stori (<http://turnoff.us/geek/>)